



OPERATING SYSTEMS

Module 2 OS Principles

C. Oswald
VIT Chennai

Networking Principles and layered architecture

- Process Concept
- Structures
 - Process Control Block
 - Ready List
 - Threads
- Operations on Processes
 - Process Creation
 - Process Termination



ADDRESSED COURSE OUTCOME

- CO-2 : Apply various types of system calls and to find the stages of various process states.

Process



- An operating system executes a variety of programs:
 - Batch system – **jobs**
 - Time-shared systems – **user programs** or **tasks**
- **Process** – a program in execution; process execution must progress in sequential fashion

Process - Contd.,

- Parts of Process
 - **Text Section:** Program code
 - Current activity including **program counter**, processor registers
 - **Stack** containing temporary data
 - Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time

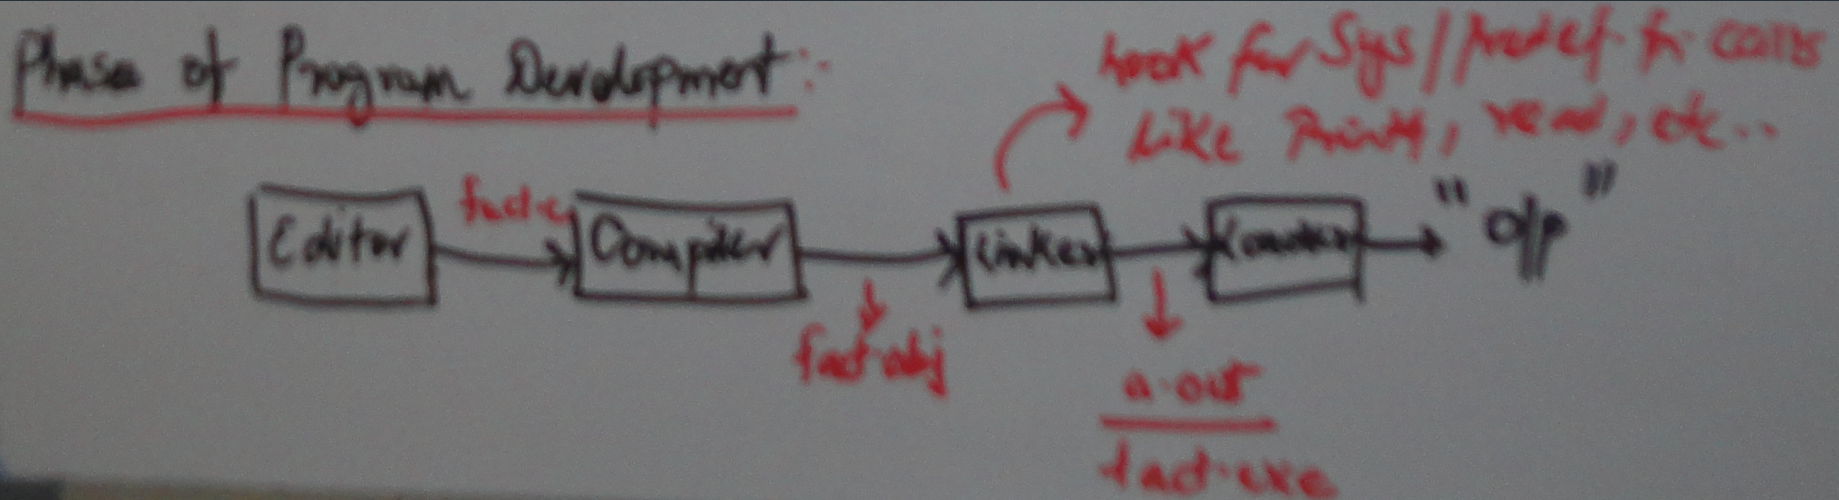
Process (Cont.)

- Program is termed as **passive** entity which stores on disk whereas process is **active** entity. Once the executable file gets loaded into memory then program will be termed as process
- Execution of program started via GUI mouse clicks, command line entry of its name, etc
- One program can be several processes
 - Consider multiple users executing the same program

PROCESS CONCEPT - Intro

➤ PROCESS NOTION

- Program in execution- Programs RUNS!!! – understand the terminology !
- Active Program (Process) v/s Binaries-Exe's on HDD (passive!)
- Program here is not Source Code!
- **Phases of PROGRAM Development (BINARY)**



Binary Development Phases

➤ Revisit of the Classical Hello World Example!

```
#include <stdio.h>
```

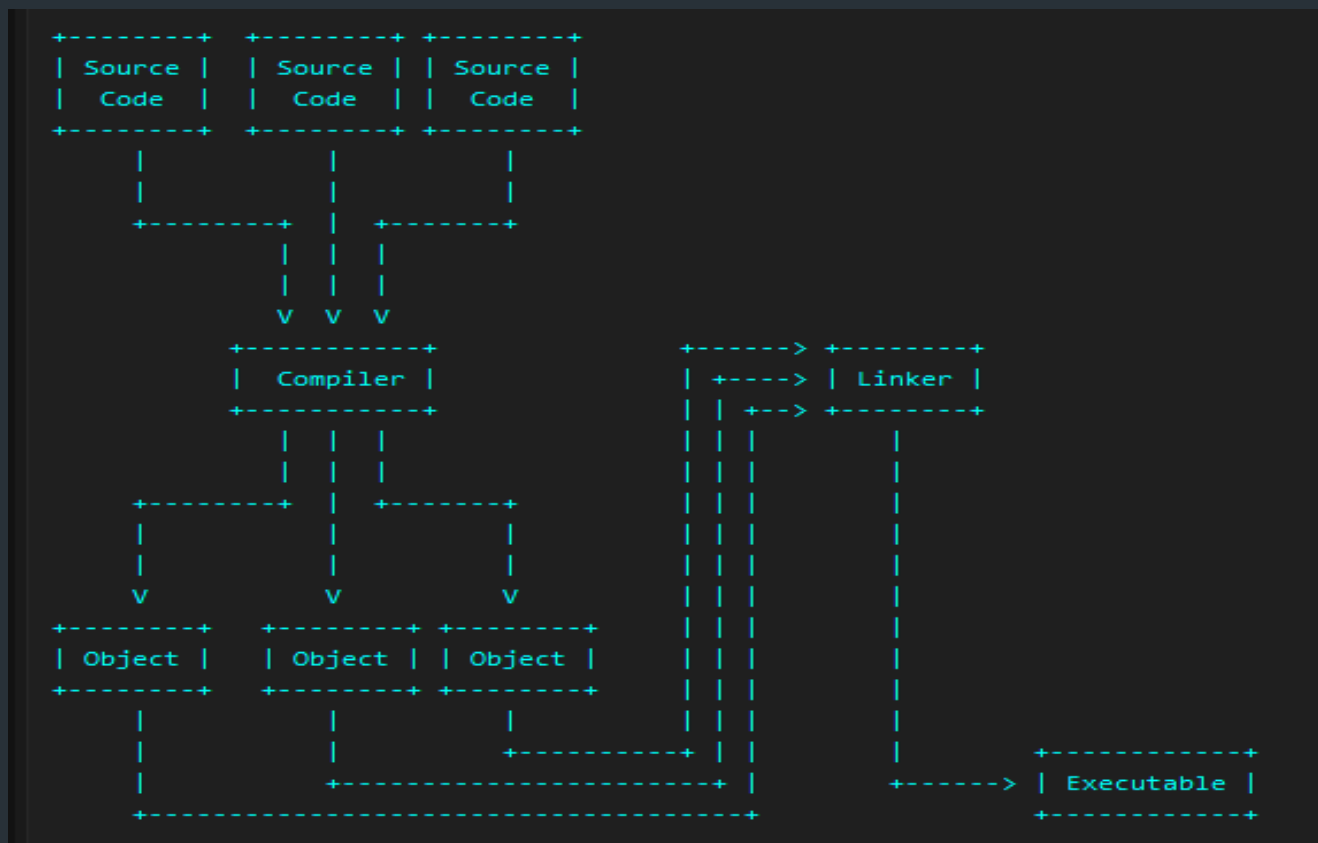
```
int main() {
```

```
printf ("Hello World \n");
```

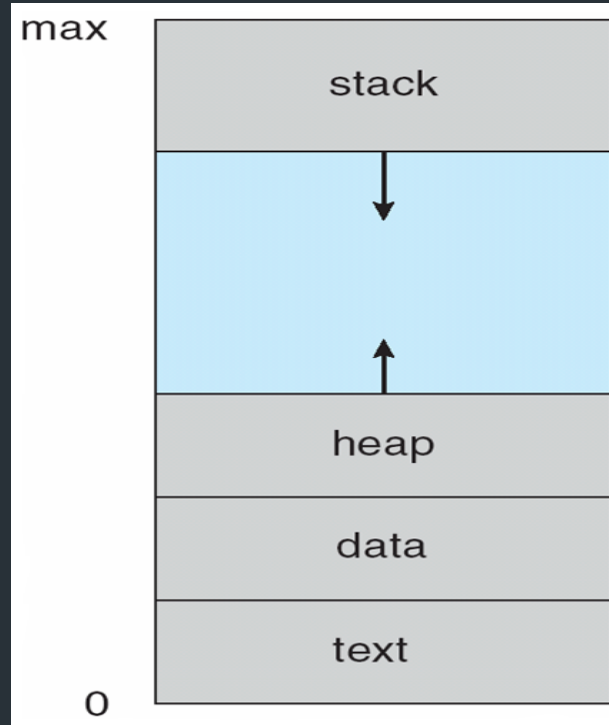
```
return 0; } /* Assume Stored in first.c */
```

- ❖ Editors in Linux – gedit, emacs, vim, etc.
- ❖ these are again programs later become processes
- ❖ Each block in the phases is a program - process
- ❖ Header File Inclusion Why ? – System v/s User Defined Headers
- ❖ What do They Contain – Function Declaration or Function Definition ; int add (int , int) ;
v/s
- ❖ int add (int x, int y) { return x+y;}
- ❖ Benefit of Having Declaration and Then Definition
- ❖ Where is Printf Defined ? - Role of Linker Justification

A Visual Understanding of the phases



Process in Memory



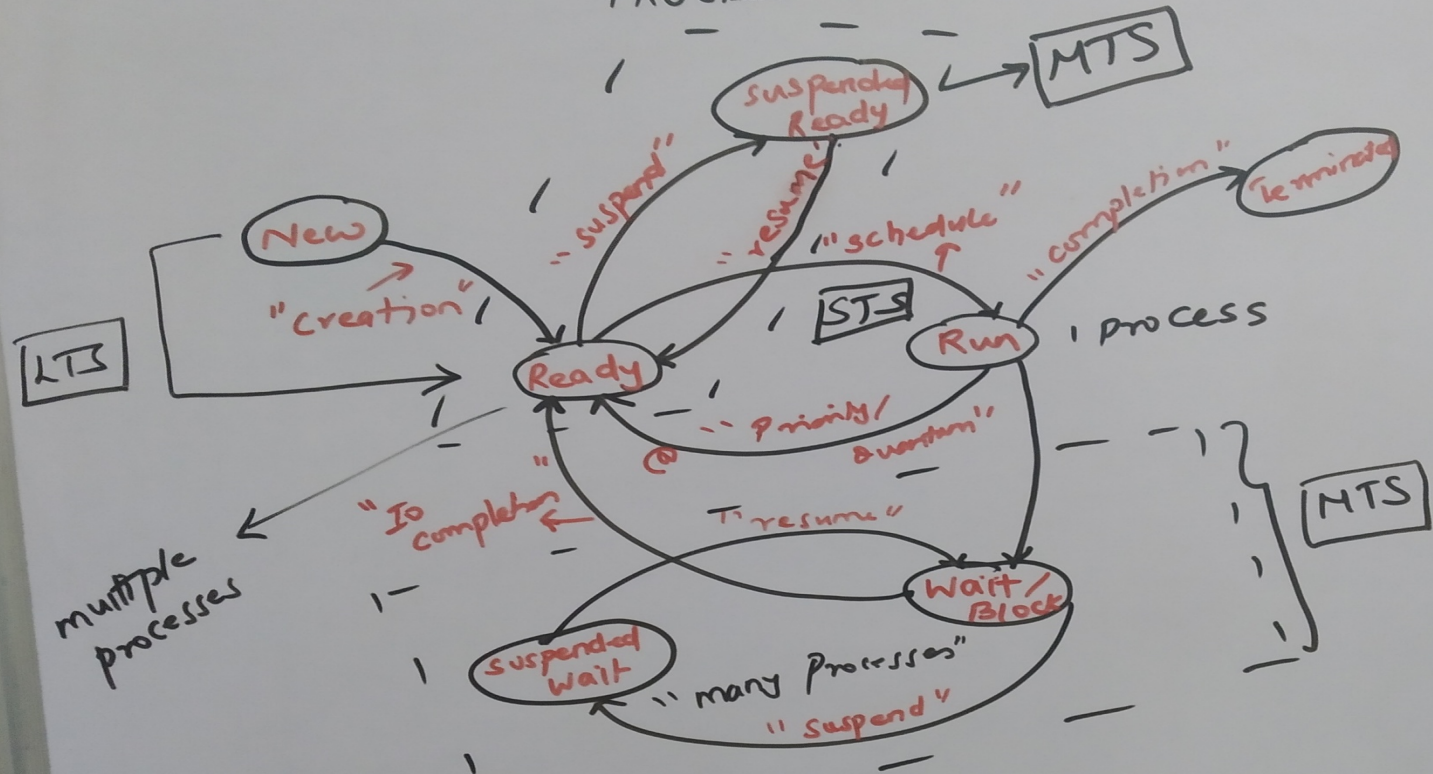
Source: Operating_System_Concepts,_8th_ Edition

Process State



- **new:** The process is being created
- **running:** Instructions are being executed
- **waiting:** The process is waiting for some event to occur
- **ready:** The process is waiting to be assigned to a processor
- **terminated:** The process has finished execution

PROCESS STATE TRANSITION



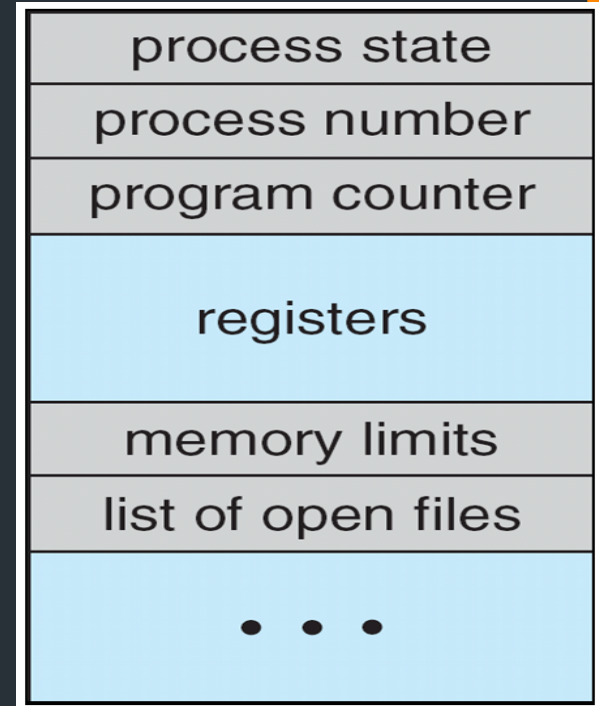
@ → Preemptive Keyway

Suspended states serve as BACKING STORE

Process Control Block (PCB)

PCB Fields:

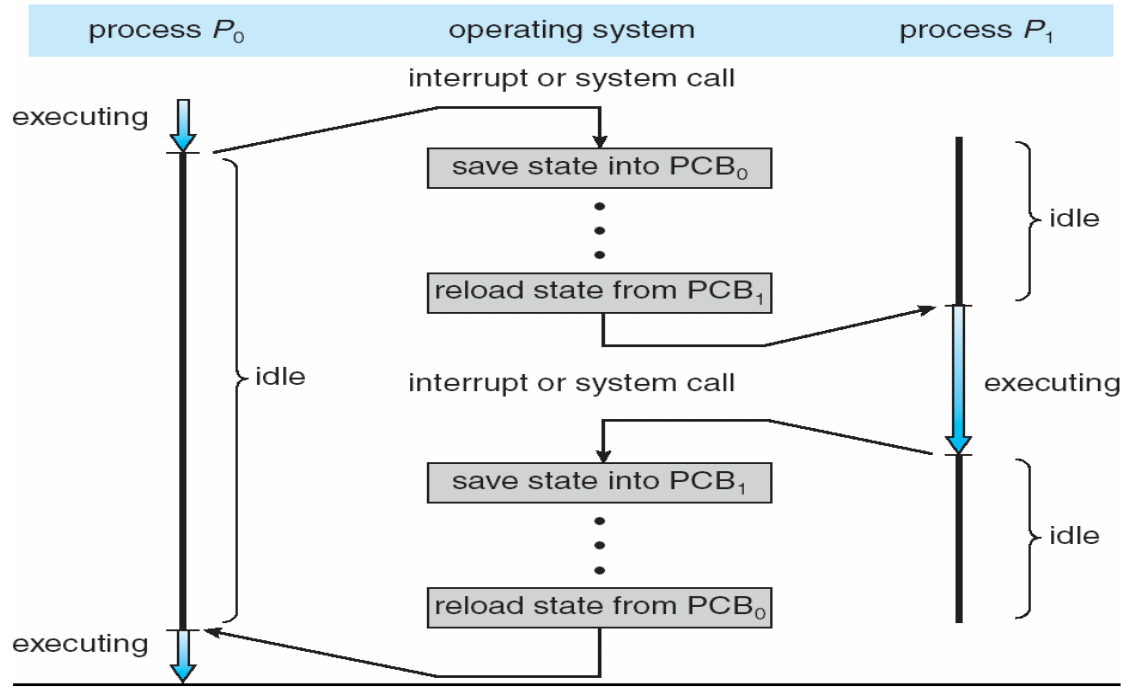
- **Process state** – new, ready, running, waiting and so on
- **Program counter** – address of the next instruction to be executed
- **CPU registers** – contents of all process-centric registers
- **CPU scheduling** information- priorities, scheduling queue pointers
- **Memory-management information** – memory allocated to the process
- **Accounting information** – CPU used, clock time elapsed since start, time limits
- **I/O status information** – I/O devices allocated to process, list of open files



Source:

Operating_System_Concepts,_8th_
Edition

CPU Switch From Process to Process



Source: Operating_System_Concepts, 8th_ Edition

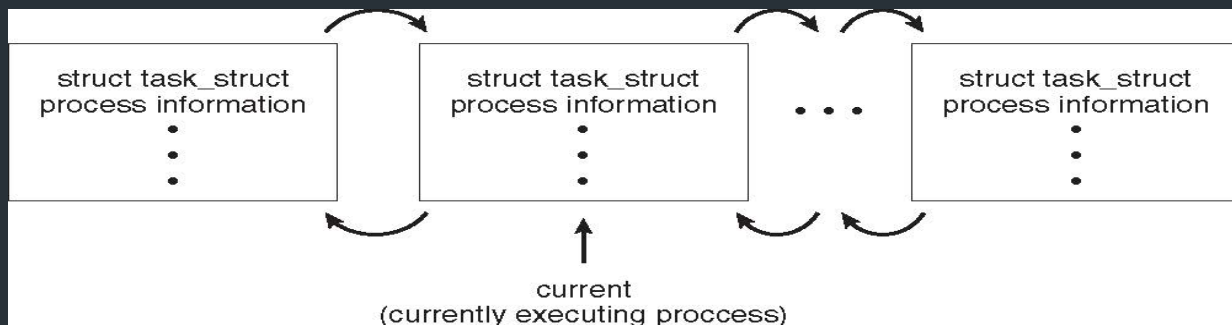
Threads

- A basic unit of CPU utilization, consisting of a **program counter, a stack, and a set of registers**, (and a **thread ID.**) Traditional (heavyweight) processes have a single **thread** of control –
 - There is one program counter, and one sequence of instructions that can be carried out at any given time.
- Consider having multiple program counters per process
 - Multiple locations can execute at once
 - Multiple threads of control -> **threads**
- Must then have storage for thread details, multiple program counters in PCB

Process Representation in Linux

Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



Process Creation



- A tree of process where Parent process create children processes, which, in turn create other processes
- process identifier (pid) : identifier used for process identification and management
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

LINUX SYSTEM CALLS OVERVIEW

System Call – Interface to OS Services (AP

✓ nothing but your C / C++ routines ; User Mode: Kernel Mode

System Call Types

Process Control ; File Manipulation ; Info Maintenance; Communication ; Protection ;
Device Management

PROCESS CONTROL Calls: `end()` ; `abort()`; `load()`; `execute()`;

`Create()`; `terminate()`; `get / set attributes ()` ; `wait()`, `fork()`, etc.

FILE MANIPULATION Calls: `create()`, `delete()`, `open()`, `read()`, `write()`, `get / set file attributes ()`;

INFORMATION MAINTENANCE calls: `getpid()`, `getppid()`, etc.

COMMUNICATION Calls: `pipe()` ; `shmopen()` ; `mmap()`, etc.

PROTECTION Calls: `chmod()` ; `chown()` ; `umask()`;

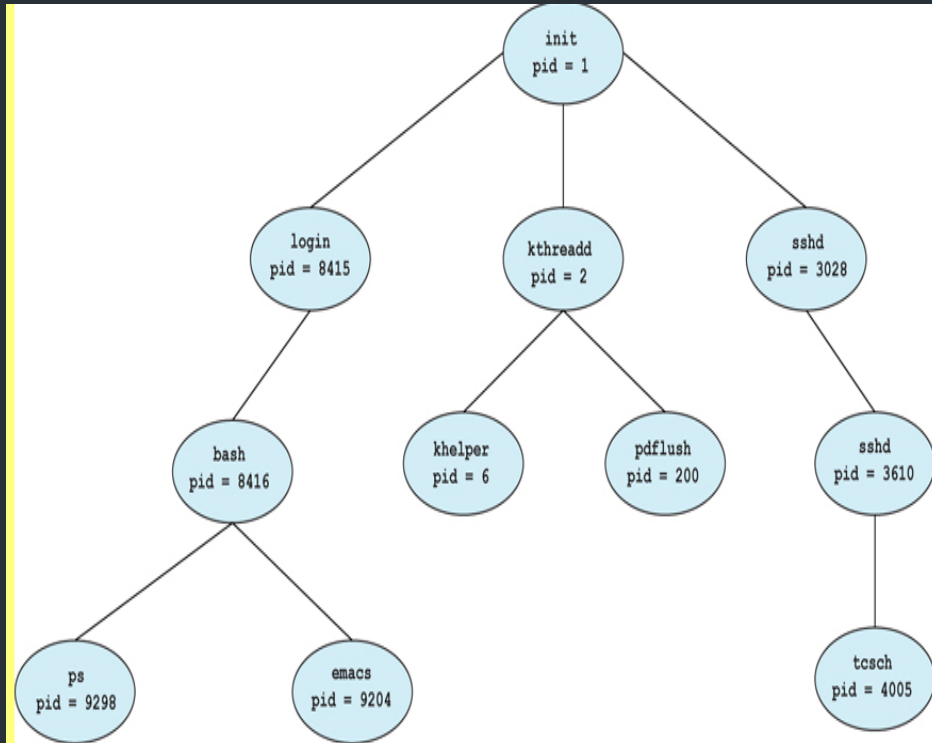
DEVICE MANAGEMENT Calls: `read()`, `write()`, `get / set device attributes`

LINUX SYSTEM CALLS OVERVIEW

FORK() and related system calls

- ✓ Process creation using fork () system call
 - ✓ Process – Program in Execution – Must reside in Main Memory, Occupy CPU
 - ✓ **Attributes of a Process**
 - ✓ PID ; STATE ; PC; PRIORITY ; GPR ;
 - ✓ List (open files); Open Devices; Protection Info
 - ✓ Process Details are stored in **Process Control Block (PCB)**
 - ✓ **Two Types of Processes** : CPU Bound v/s IO Bound
 - ✓ Running State v/s Wait State
- [More Cpu Time v/s More IO Time] –
Right Balance of processes – Schedulers Challenge

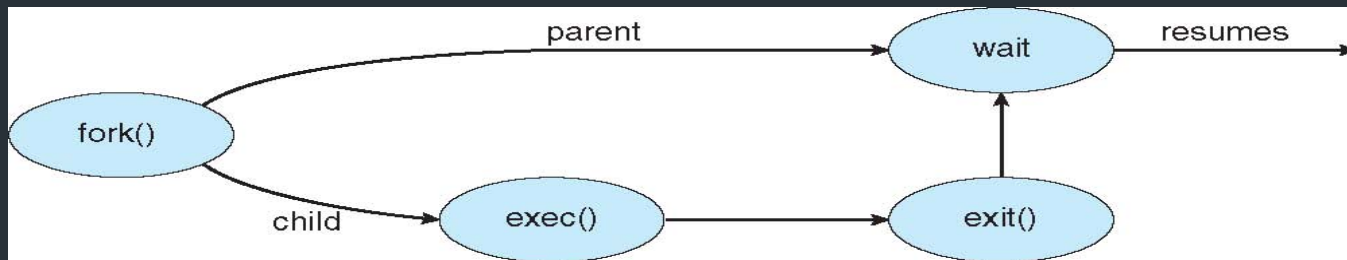
Sample Tree of Processes in Linux



- On typical LINUX systems the process scheduler is termed **sched**, and is given PID 0. The first thing it does at system startup time is to launch **init**, which gives that process PID 1. Init then launches all system daemons and user logins, and becomes the ultimate parent of all other processes.

Process Creation (Cont.)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - `fork()` system call creates new process
 - `exec()` system call used after a `fork()` to replace the process' memory space with a new program




PROCESS MGMT CONTINUED

❑ 3 types of schedulers in Linux – LTS ; STS ; MTS

- ✓ LTS – New Process to Ready State (also called as Job Scheduler)
- ✓ STS – Ready to Run (also called as CPU Scheduler)
- ✓ MTS – Suspend to Resume (vice versa) – Swapper
- ✓ Feeling of Seamless or Infinite Process Creation for End User

Despite internal limit on no of processes

- ❖ Degree of Multiprogramming – No of Processes in Main Memory at a given time t – controlled by the LTS
- ❖ LTS Challenge – right balance of CPU bound and IO Bound processes – achieve good throughput
- ❖ Swapper – serves the purpose of Context Switching

- 
- ✓ cc first.c
 - ✓ ./a.out and enter – this is parent process in execution
 - ✓ Leads a call to main – at fork point – a copy of a.out is created and this is the child process
 - ✓ By default child carries the same image as that of the parent post the forking point
 - ✓ Order of execution of parent / child operations / statements is in kernel's schedulers hands!
 - ✓ Next printf example – output statements ordering can change.
 - ✓ Assume parent gets the control first (which again can vary)
 - ✓ Subsequent sessions – other calls such as exec, variants of it, wait, etc.
 - ✓ Overlapped with Scheduling Algorithms – Preemptive v/s Non Preemptive

FORK System Call

❖ Why is it called Fork() – dictionary relevance!

- ✓ Helps in creating New Processes in Linux
- ✓ New Process is created as child of the parent process
- ✓ Parent is the calling process (normally ./a.out or main call)
- ✓ On successful fork – Two processes reside in Main Memory (Parent and Child)

❖ Child Process definition – statements post the forking point


❖ Three Possible Return values of Fork() Sytem call

0 indicates child process control

>0 indicates parent process control

<0 failed fork call

Content Switching



CST → Content Switch Time

— X —
FIRST FORK EXAMPLE

```
int main() {
```

```
    int pid;
```

```
    pid = fork();
```

```
    if (pid < 0)
```

```
        printf("Fork Failed \n");
```

```
    else if (pid == 0)
```

```
        printf("child Block \n");
```

```
    else if (pid > 0)
```

```
        printf("Parent Block \n");
```


```
    return 0;
```

```
}
```

Parent Block
child Block

0/p

Content Switching



CST → Content Switch Time

FIRST FORK EXAMPLE

```
int main() {
```

```
    int pid;
```

```
    pid = fork();
```

How many times
will "Welcome to os"
print?

```
    if (pid < 0)
```

```
        printf("Fork Failed \n");
```

```
    else if (pid == 0)
```

```
        printf("child Block \n");
```

```
    else if (pid > 0)
```

```
        printf("Parent Block \n");
```

```
        printf("Welcome To os \n");
```

```
    return 0;
```

C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

Source: Operating_System_Concepts,_8th_ Edition

Process Termination

- Process executes last statement and then asks the operating system to delete it using the `exit()` system call.
- Parent may terminate the execution of children processes using the `abort()` system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Process Termination

- If a process terminates, then all its children must also be terminated.
 - **cascading termination.** All children, grandchildren, etc. are terminated.
 - The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```

- If no parent waiting (did not invoke `wait()`) process is a **zombie**
- If parent terminated without invoking `wait`, process is an **orphan**

Multiprocess Architecture – Chrome Browser

- Many web browsers ran as single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
 - **Browser** process manages user interface, disk and network I/O
 - **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in





References

- Abraham Silberschatz, Peter B. Galvin, Greg Gagne-Operating System Concepts, Wiley (2018).



■ **THANK YOU**