

Module 5 – L2:

Memory management

Paging - Demand Paging - Page Faults

Dr. Rishikeshan C A
Assistant Professor (Sr.)
SCOPE, VIT Chennai

Outline

- **Paging**
- **Address Translation Scheme**
- **Implementing paging, Paging H/W**
- **Paging Hardware With TLB**
- **Effective Access Time (EAT)**
- **Demand Paging**
- **Page Faults**
- **Page Fault Handling**

Background


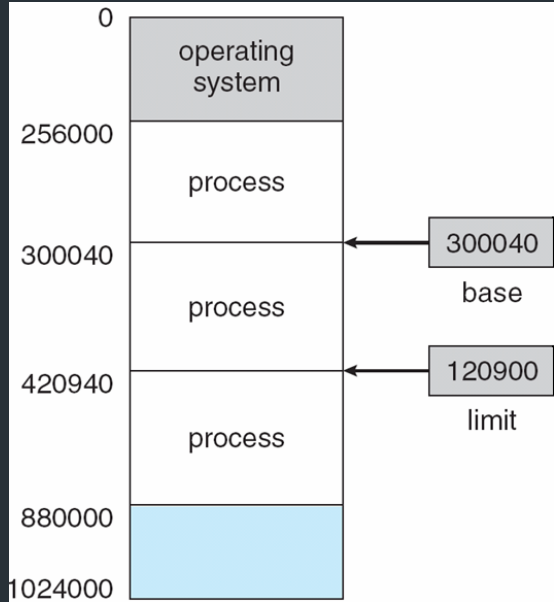


- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are only storage CPU can access directly
- Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- Register access in one CPU clock (or less)
- Main memory access can take many cycles
- **Cache** sits between main memory and CPU registers
- Protection of memory required to ensure correct operation

Base and Limit Registers

4

- A pair of **base** and **limit registers** define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user



Hardware Address Protection

Ref:

Logical vs. Physical Address Space

5

- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program


Memory-Management Unit (MMU)

6

- Hardware device that at run time maps virtual to physical address
- To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
 - Base register now called **relocation register**
 - MS-DOS on Intel 80x86 used 4 relocation registers
- The user program deals with *logical* addresses; it never sees the *real* physical addresses
 - Execution-time binding occurs when reference is made to location in memory
 - Logical address bound to physical addresses

Dynamic relocation using a relocation register 7


- ❑ Better memory-space utilization; unused routine is never loaded
- ❑ All routines kept on disk in relocatable load format
- ❑ Useful when large amounts of code are needed to handle infrequently occurring cases
- ❑ No special support from the operating system is required
 - ❑ Implemented through program design
 - ❑ OS can help by providing libraries to implement dynamic loading



- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
 - Total physical memory space of processes can exceed physical memory
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk

Schematic View of Swapping

9



Swapping (Cont.)

10

- Does the swapped out process need to swap back in to same physical addresses?
 - Depends on address binding method
 - Plus consider pending I/O to / from process memory space
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - Swapping normally disabled. Started if more than threshold amount of memory allocated
 - Disabled again once memory demand reduced below threshold

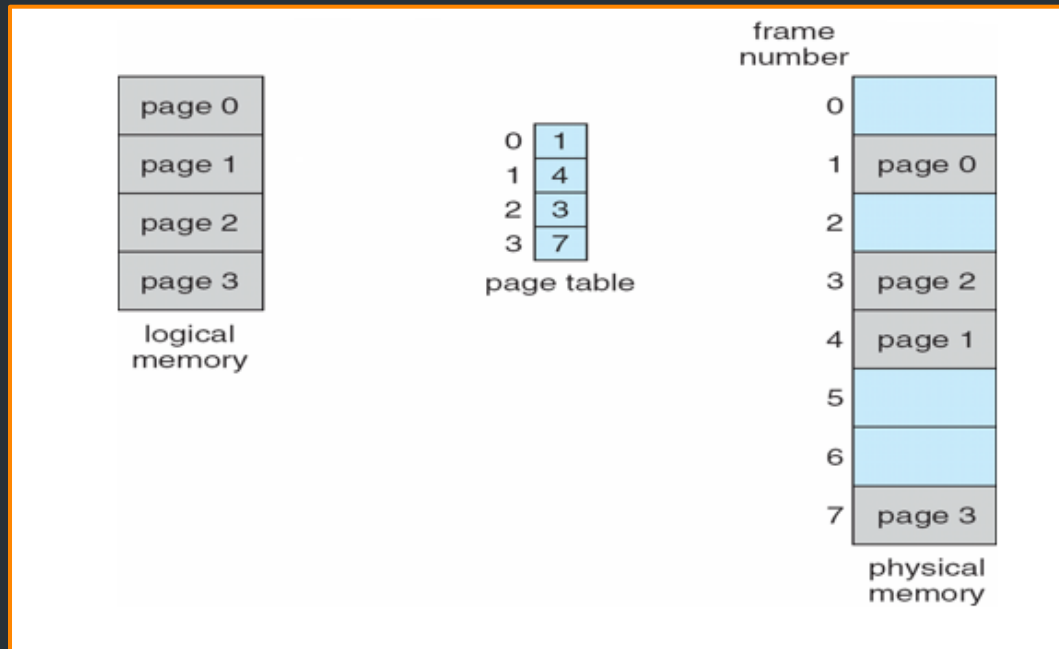
Paging

11

- Idea: Divide memory into fixed-sized **pages**
- Divide logical memory into blocks of same size **pages**.
 - Both virtual and physical memory are composed of pages
 - Divide physical memory into fixed-sized chunks/blocks called **frames** (size is power of 2, usually between 512 bytes and 16 MB).
- Goal
 - Eliminate external fragmentation
 - Don't allocate memory that will not be used

Paging: Example

12



Paging Implementation

- Keep track of all free physical page frames
- To run a program of size n pages, need to find n free frames and load program
 - Pages **do not** need to be contiguous!
- Set up mapping from virtual to physical in **page table**
- At memory reference time, translate virtual address to physical address using **page table**.

Paging: free-frame list

13



Before allocation

After allocation

- To run a program of size n pages, need to find any n free frames and load all the program (pages).
- So need to keep track of all free frames in physical memory use free-frame list.

Logical address used in paging

14

- Within each program, each logical address must consist of a page number and an offset within the page.
- A dedicated register always holds the starting physical address of the page table of the currently running process.
- Presented with the logical address (page number, offset) the processor accesses the page table to obtain the physical address (frame number, offset).

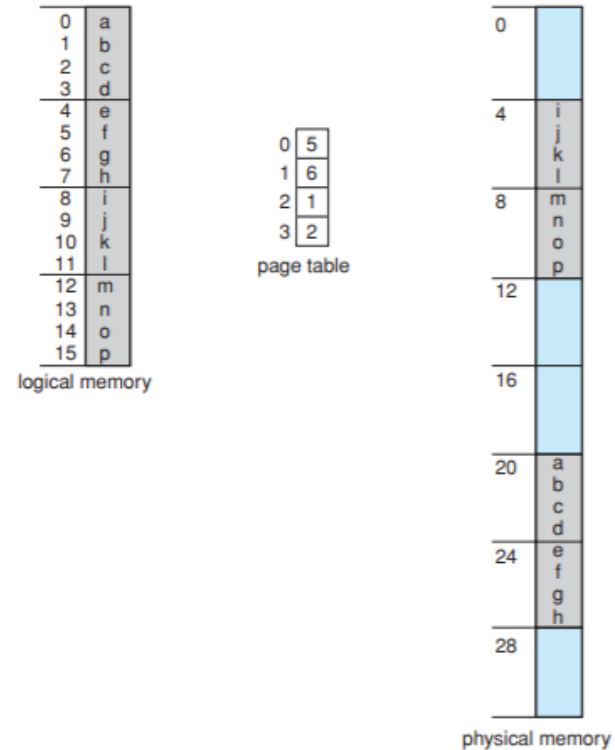



Figure: Paging example for a 32-byte memory with 4-byte pages

Address Translation Scheme

15

- Address generated by CPU is divided into:
 - Page number (p)** – used as an index into a *page table* which contains base address of each page in physical memory
 - Page offset/displacement (d)** – combined with base address to define the physical memory address that is sent to the memory unit




- For given logical address space 2^m and page size 2^n

Example:

- 4 KB (=4096 byte) pages
- 32 bit logical addresses

$$2^d = 4096 \longrightarrow d = 12$$

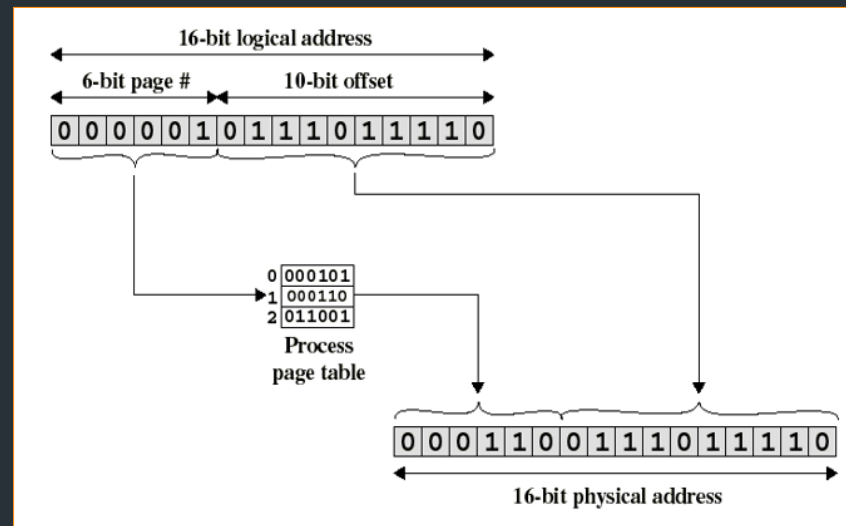


Logical-to-Physical Address Translation in Paging

16

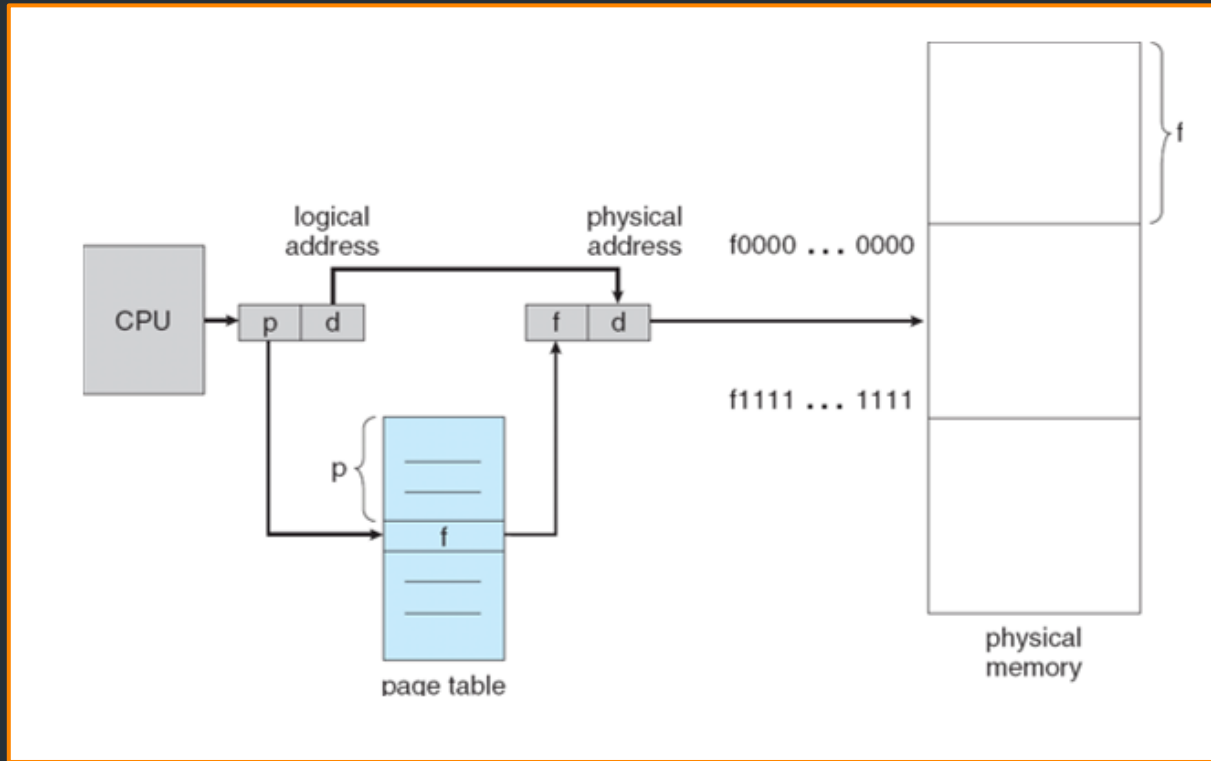
By using a page size of a power of 2, the pages are invisible to the programmer, compiler/assembler, and the linker.

Address translation at run-time is then easy to implement in hardware: logical address (p, d) gets translated to physical address (f, d) by indexing the page table with p and appending the same displacement/offset d to the frame number f.



Address Translation Architecture (Paging Hardware)

17



Internal Fragmentation & Frame size

- Page size = 2,048 bytes
- Process size = 72,766 bytes
- 35 pages + 1,086 bytes
- Internal fragmentation of $2,048 - 1,086 = 962$ bytes
- Worst case fragmentation = 1 frame – 1 byte
- On average fragmentation = $1 / 2$ frame size
- So small frame sizes desirable?
- But each page table entry takes memory to track.
- Page sizes growing over time:
 - Solaris supports two page sizes – 8 KB and 4 MB.

Two-Level Paging Example

19



- A logical address (on 32-bit machine with 1K page size) is divided into:
 - a page number consisting of 22 bits
 - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
 - a 12-bit page number
 - a 10-bit page offset

where p_i is an index into the outer page table, and p_2 is the displacement within the page of the outer page table

- Thus, a logical address is as follows:

page number		page offset
p_i	p_2	d
12	10	10

Example


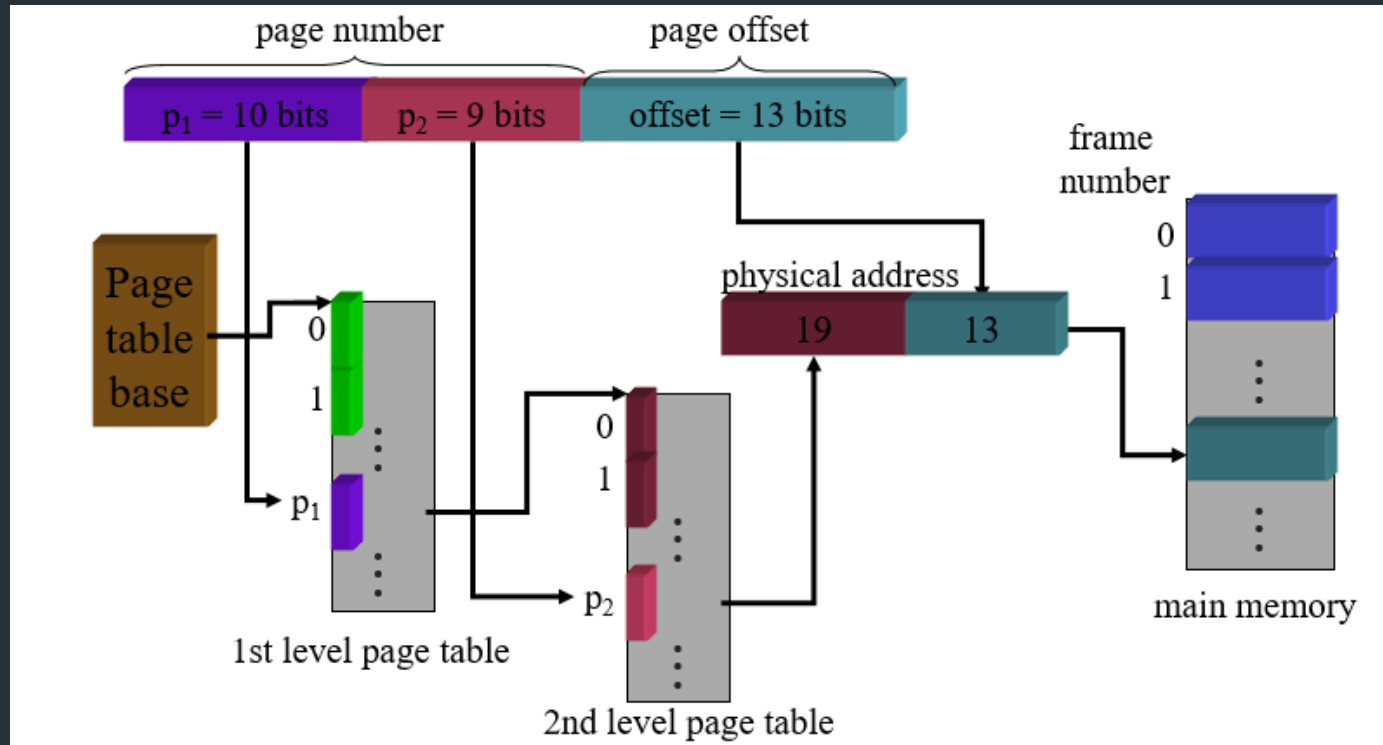


Illustration of Address-Translation Scheme

20

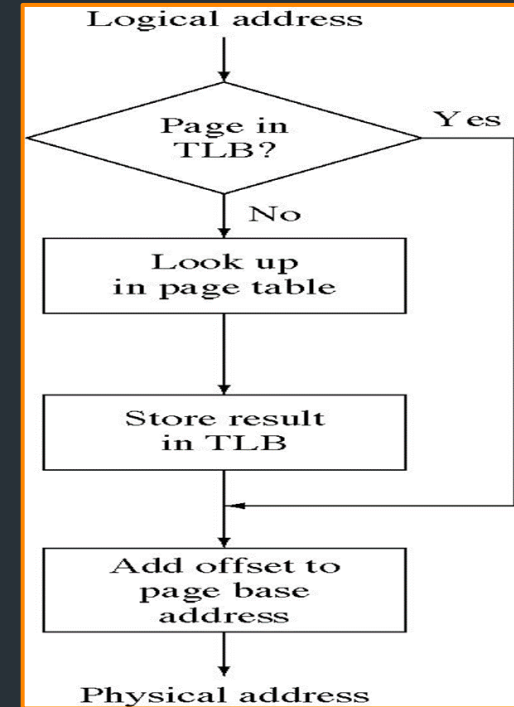



Implementing page tables in hardware

- Page table resides in main (physical) memory
- CPU uses special registers for paging
 - Page table base register (PTBR) points to the page table
 - Page table length register (PTLR) contains length of page table: restricts maximum legal logical address
- Translating an address requires two memory accesses
 - First access reads page table entry (PTE)
 - Second access reads the data / instruction from memory
- Reduce number of memory accesses
 - Can't avoid second access (we need the value from memory)
 - Eliminate first access by keeping a hardware cache *Associative Memory (Registers)* (called a *translation lookaside buffer* or TLB) of recently used page table entries: enables fast parallel search

Paging Hardware With TLB

22



How TLB works

- TLB takes advantage of the **Locality Principle**.
- TLB uses associative mapping hardware to simultaneously interrogate all TLB entries to find a match/hit on page number.
- TLB hit rates are 90+%.
- The TLB must be flushed each time a new process enters the running state.
- Maybe keep/load TLB information in/from process context.

Effective Access Time (EAT)

- Effective Access Time (EAT) is between 1 and 2 access times – should be closer to 1.
- Assume memory cycle time is 1 microsecond.
- Associative (Memory) Lookup = ε time unit.
- Hit ratio = α – percentage of times that a page number is found in the associative memory; ratio related to number of associative registers.
- $$\text{EAT} = \alpha(\varepsilon + 1) + (1 - \alpha)(\varepsilon + 2) = 2 + \varepsilon - \alpha$$

EAT calculation example

- Assume memory cycle time is 100 nanosecond.
- Associative memory lookup = 20 nanosecond.
- Hit ratio = 80%
 - $EAT = 0.80 \times 120 + 0.20 \times 220 = 140 \text{ ns}$
 - So 40% slowdown in memory access time.
- Hit ratio = 98%
 - $EAT = 0.98 \times 120 + 0.02 \times 220 = 122 \text{ ns}$
 - So only 22% slowdown in memory access time.

Page Size

26

Small page size

- Advantages
 - less internal fragmentation
 - better fit for various data structures, code sections
 - less unused program in memory
- Disadvantages
 - programs need many pages, larger page tables

Overhead due to page table and internal fragmentation

The diagram shows the formula for overhead: $overhead = \frac{s \cdot e}{p} + \frac{p}{2}$. The first term, $\frac{s \cdot e}{p}$, is enclosed in a circle and labeled 'page table space' with an arrow. The second term, $\frac{p}{2}$, is also enclosed in a circle and labeled 'internal fragmentation' with an arrow.

Optimized when
 $p = \sqrt{2se}$

Where,
s = average process size in bytes
p = page size in bytes
e = page entry

Demand paging


27

- **Virtual memory** – separation of user logical memory from physical memory
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
 - Allows for more efficient process creation
 - More programs running concurrently
 - Less I/O needed to load or swap processes
- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation

Demand Paging

28

- Bring a page into memory only when it is needed.
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- Page is needed \Rightarrow reference to it
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory




Virtual Memory That is Larger Than Physical Memory

Demand Paging

29


- Similar to paging system with swapping (diagram on right)
- **Lazy swapper** – never swaps a page into memory unless page will be needed
 - Swapper that deals with pages is a **pager**
- If pages needed are already **memory resident**
 - No difference from non demand-paging
- If page needed and not memory resident
 - Need to detect and load the page into memory from storage
 - Without changing program behavior
 - Without programmer needing to change code



Valid-Invalid Bit

30

- With each page table entry a valid–invalid bit is associated
(**v** \Rightarrow in-memory – **memory resident**,
i \Rightarrow not-in-memory)
- Initially valid–invalid bit is set to **i** on all entries
- During MMU address translation, if valid–invalid bit in page table entry is **i** \Rightarrow page fault




Page table when some pages are not in main memory

Page Fault & Page fault handling

1. If there is a reference to a page, first reference to that page will trap to operating system: **Page fault**
2. Operating system looks at another table to decide:
 - Invalid reference \Rightarrow abort
 - Just not in memory
3. Find free frame
4. Swap page into frame via scheduled disk operation
5. Reset tables to indicate page now in memory set validation bit = **v**
6. Restart the instruction that caused the page fault

Steps in Handling a Page Fault

32



What happens if there is no free frame?

33

- Page replacement – find some page in memory, but not really in use, swap it out.
 - page replacement algorithms
 - performance – want an algorithm which will result in minimum number of page faults.
- Same page may be brought into memory several times.

Performance of Demand Paging

- Page Fault Rate $0 \leq p \leq 1.0$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault
- Effective Access Time (EAT)
$$\text{EAT} = (1 - p) \times \text{memory access}$$
 - + p (page fault overhead
 - + [swap page out]
 - + swap page in
 - + restart overhead)

Demand Paging Example

- Suppose that a normal memory access requires 200 nanoseconds, and that servicing a page fault takes 8 milliseconds (8,000,000 nanoseconds, or 40,000 times a normal memory access) with a page fault rate of p , (on a scale from 0 to 1), the effective access time (EAT) is now:

$$(1 - p) * (200) + p * 8000000$$

$$= 200 + 7,999,800 * p$$

- So, EAT clearly depends heavily on p ! Even if only one access in 1000 causes a page fault, the effective access time drops from 200 nanoseconds to 8.2 microseconds, a slowdown of a factor of 40 times.
- In order to keep the slowdown less than 10%, the page fault rate must be less than 0.0000025, or one in 399,990 accesses.

Aspects of Demand Paging

35

- Extreme case – start process with *no* pages in memory
 - OS sets instruction pointer to first instruction of process, non-memory-resident -> page fault
 - And for every other process pages on first access
 - Pure demand paging
- Actually, a given instruction could access multiple pages -> multiple page faults
 - Consider fetch and decode of instruction which adds 2 numbers from memory and stores result back to memory
 - Pain decreased because of *locality of reference*
- Hardware support needed for demand paging
 - Page table with valid / invalid bit
 - Secondary memory (swap device with *swap space*)
 - Instruction restart



Any Questions ?

References

- <http://www.cs.nthu.edu.tw/~ychung/slides/CSC3150/Abraham-Silberschatz-Operating-System-Concepts---9th2012.12.pdf>
- <http://www.digilife.be/quickreferences/qrc/linux%20system%20call%20quick%20reference.pdf>
- [Silberschatz, Gagne, Galvin: Operating System Concepts, 6th Edition](#)

Thank You!