



Module 5:

Memory management

Segmentation, Page Replacement algorithms,
Thrashing , Working Set

Dr. Rishikeshan C A
Assistant Professor (Sr.)
SCOPE, VIT Chennai

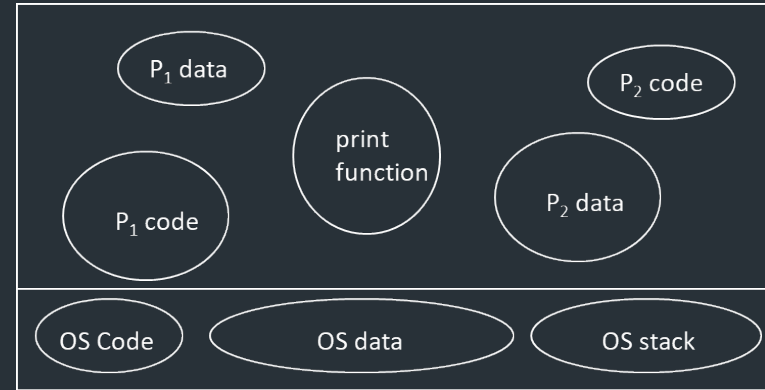
Outline

- Segmentation
- Segmentation Hardware
- Segmentation Architecture
- Pros and Cons of Segmentation
- Page-Replacement Algorithms
- Thrashing
- Working Set

Segmentation

3

- **Segment**: a region of logically contiguous memory
- **Segmentation-based transition**: use a table of base-and-bound pairs
- Memory-management scheme that supports user view of memory
- A program is a collection of segments. A segment is a logical unit such as:
 - main program,
 - procedure,
 - function,
 - method,
 - object,
 - local variables, global variables,
 - common block,
 - stack,
 - symbol table, arrays




logical address space

Segmentation

4

- Segmentation is a technique for breaking memory up into logical pieces
- Each “piece” is a grouping of related information
 - data segments for each process
 - code segments for each process
 - data segments for the OS
 - etc.
- Like paging, use virtual addresses and use disk to make memory look bigger than it really is
- Segmentation can be implemented with or without paging



User's View of a Program

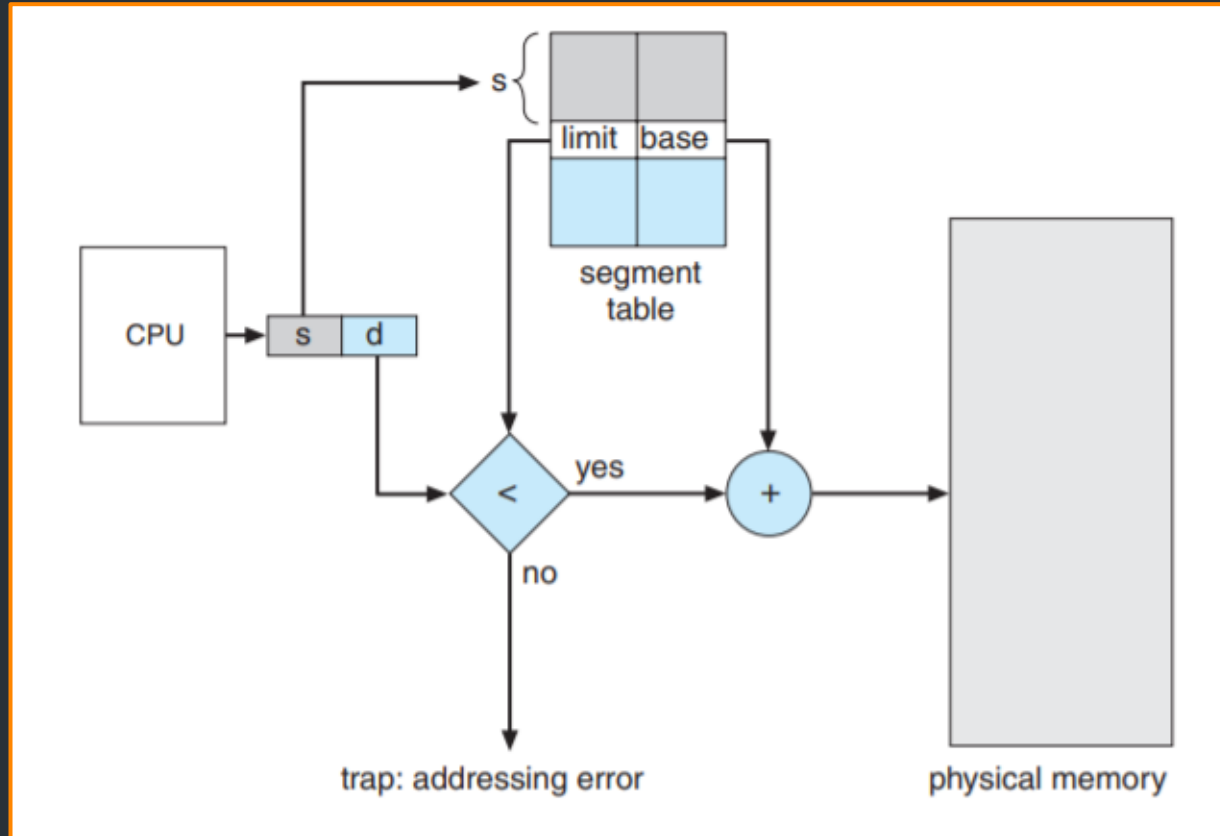
Addressing Segments

5

- Let's first assume no paging in the system
- User generates logical addresses
- These addresses consist of a segment number and an offset into the segment
- Use segment number to index into a table
- Table contains the physical address of the start of the segment
 - often called the base address
- Add the offset to the base and generate the physical address
 - before doing this, check the offset against a limit
 - the limit is the size of the segment

Addressing Segments

6



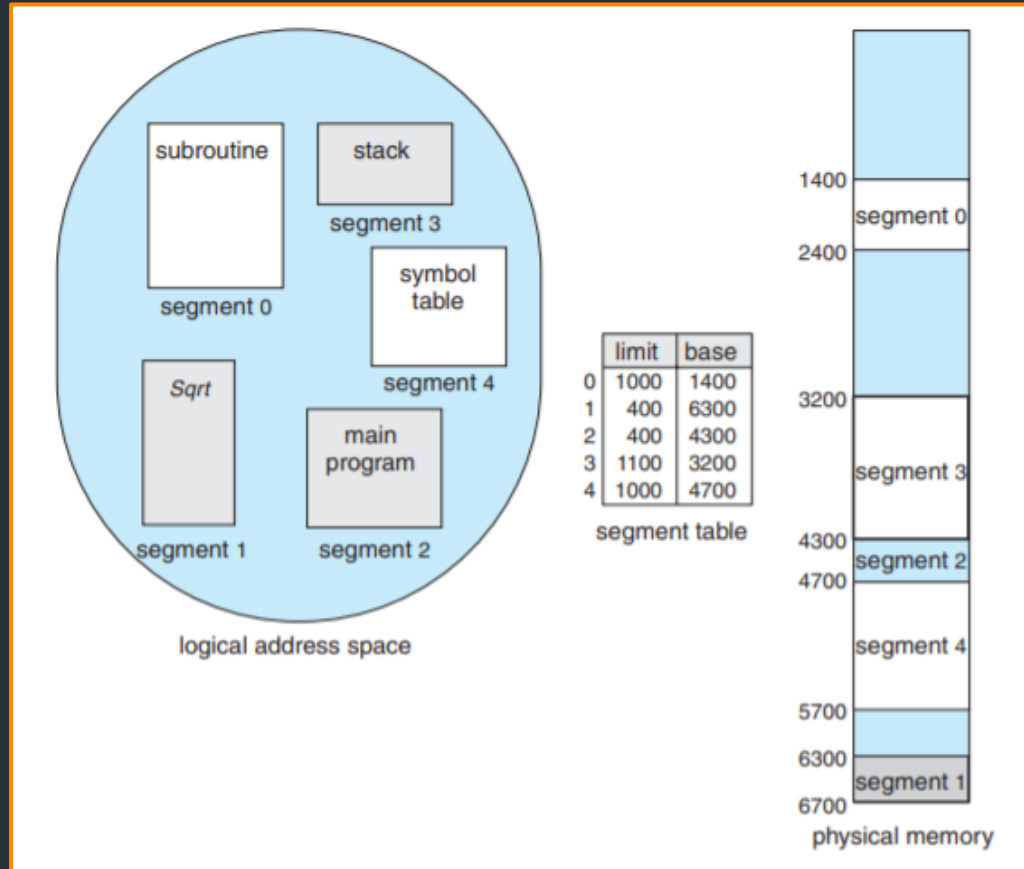
Segmentation Hardware

7

- Sounds very similar to paging
- Big difference – segments can be variable in size
- As with paging, to be effective hardware must be used to translate logical address
- Most systems provide segment registers
- If a reference isn't found in one of the segment registers
 - trap to operating system
 - OS does lookup in segment table and loads new segment descriptor into the register
 - return control to the user and resume
- Again, similar to paging

Example of Segmentation

8



Segmentation Architecture

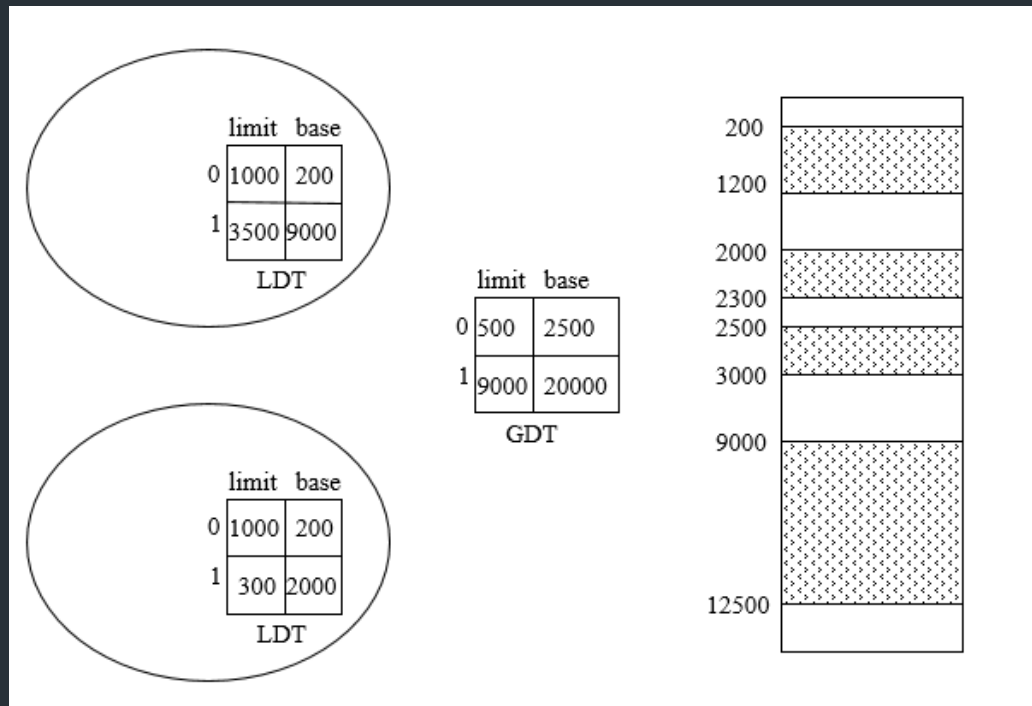
- Logical address consists of a two tuple:
 <segment-number, offset>,
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - **base** – contains the starting physical address where the segments reside in memory
 - **limit** – specifies the length of the segment
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;
 segment number **s** is legal if **s** < **STLR**

Segmentation Architecture (Cont.)

- Protection
 - With each entry in segment table associate:
 - validation bit = 0 \Rightarrow illegal segment
 - read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- A segmentation example is shown in the following diagram

Protection and Sharing

- Like page tables, each process usually gets its own segment table
- Unlike page tables, there usually exists a global segment table for everyone
 - this, however, is usually used by OS
- Access rights for segment are usually included in table entry
- Multiple processes can share a segment



Pros and Cons of Segmentation

- Easier to grow and shrink individual segments
- Finer control of segment accesses
 - e.g., read-only for shared code segment
- More efficient use of physical space
- Multiple processes can share the same code segment
- Memory allocation is still complex
 - Requires contiguous allocation
- Entire segment is either in memory or on disk
- Variable sized segments leads to external fragmentation in memory
- Must find a space big enough to place segment into
- May need to swap out some segments to bring a new segment in

Segmentation with Paging


- Most architectures support segmentation and paging
- Basic idea,
 - segments exist in virtual address space
 - base address in segment descriptor table is a virtual address
 - use paging mechanism to translate this virtual address into a physical address
- Now an entire segment does not have to be in memory at one time
 - only the part of the segment that we need will be in memory



Page Replacement Algorithms

Need For Page Replacement

15



Basic Page Replacement


16

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim frame**
 - Write victim frame to disk if dirty
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Continue the process by restarting the instruction that caused the trap

Note: potentially 2 page transfers for page fault – increasing EAT

Page Replacement mechanism

17



Page Replacement Algorithms

18

- **Page-replacement algorithm**
 - Want lowest page-fault rate on both first access and re-access
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
 - String is just page numbers, not full addresses
 - Repeated access to the same page does not cause a page fault
 - Results depend on number of frames available
- In all our examples, the **reference string** of referenced page numbers is

7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1

First-In-First-Out (FIFO) Algorithm

19

- Reference string: **7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1**
- 3 frames (3 pages can be in memory at a time per process)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1


7	7	7	2	2	2	4	4	4	0	0	0	7	7	7
	0	0	0	3	3	3	2	2	2	1	1	1	0	0
		1	1	1	0	0	0	3	3	3	2	2	2	1

page frames

15 page faults

- Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
 - Adding more frames can cause more page faults!
- ▶ **Belady's Anomaly**
- How to track ages of pages?
 - Just use a FIFO queue

Illustrating Belady's Anomaly




Optimal Algorithm

20

- Replace page that will not be used for longest period of time
 - 9 is optimal for the example
- How do you know this?
 - Can't read the future
- Used for measuring how well your algorithm performs

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

9 page faults


Least Recently Used (LRU) Algorithm

21

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used
- But how to implement?

LRU Algorithm (Cont.)

22

- Counter implementation
 - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - When a page needs to be changed, look at the counters to find smallest value
 - Search through table needed
- Stack implementation
 - Keep a stack of page numbers in a double link form:
 - Page referenced:
 - move it to the top
 - requires 6 pointers to be changed
 - But each update more expensive
 - No search for replacement
- LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly

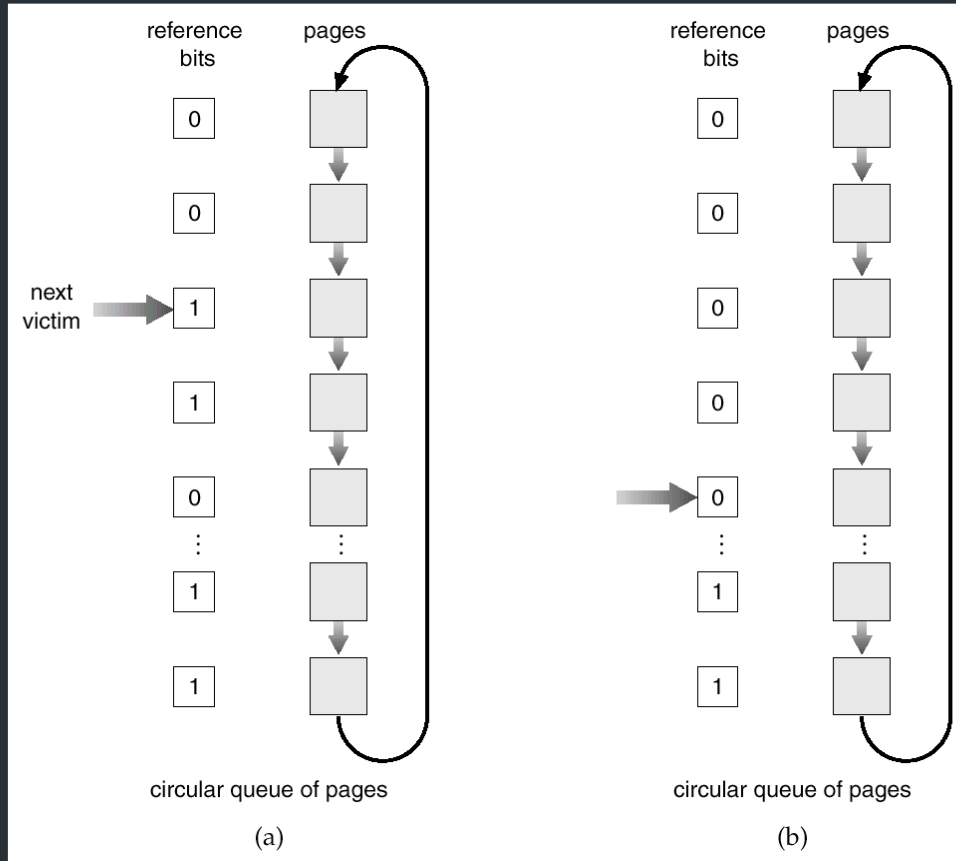
Clock page replacement algorithm (LRU Approximation Algorithms)

23

- Reference bit
 - With each page associate a bit, initially = 0
 - When page is referenced bit set to 1.
 - Replace the one which is 0 (if one exists). We do not know the order, however.
- Second chance
 - Need reference bit.
 - Clock replacement.
 - If page to be replaced (in clock order) has reference bit = 1. then:
 - set reference bit 0.
 - leave page in memory.
 - replace next page (in clock order), subject to same rules.

Second-Chance (clock) Page-Replacement Algorithm

24




Counting Algorithms

- Keep a counter of the number of references that have been made to each page
 - Not common
- **Least Frequently Used (LFU) Algorithm:** replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm:** based on the argument that the page with the smallest count was probably just brought in and has yet to be used

- If a process does not have “enough” pages, the page-fault rate is very high
 - Page fault to get page
 - Replace existing frame
 - But quickly need replaced frame back
 - This leads to:
 - Low CPU utilization
 - Operating system thinking that it needs to increase the degree of multiprogramming
 - Another process added to the system
- **Thrashing** \equiv a process is busy swapping pages in and out

Thrashing

27




Source: Operating System Concepts, 9th Edition, **Abraham Silberschatz**, Peter B. Galvin, Greg Gagne

Working-Set Model

28

- $\Delta \equiv$ working-set window \equiv a fixed number of page references
Example: 10,000 instructions
- WSS_i (working set of Process P_i) = total number of pages referenced in the most recent Δ (varies in time)
 - if Δ too small will not encompass entire locality
 - if Δ too large will encompass several localities
 - if $\Delta = \infty \Rightarrow$ will encompass entire program
- $D = \sum WSS_i \equiv$ total demand frames
 - Approximation of locality
- if $D > m \Rightarrow$ Thrashing
- Policy if $D > m$, then suspend or swap out one of the processes



Keeping Track of the Working Set


29

- Approximate with interval timer + a reference bit
- Example: $\Delta = 10,000$
 - Timer interrupts after every 5000 time units
 - Keep in memory 2 bits for each page
 - Whenever a timer interrupts copy and sets the values of all reference bits to 0
 - If one of the bits in memory = 1 \Rightarrow page in working set
- Why is this not completely accurate?
- Improvement = 10 bits and interrupt every 1000 time units

Page-Fault Frequency

30


- More direct approach than WSS
- Establish “acceptable” **page-fault frequency (PFF)** rate and use local replacement policy
 - If actual rate too low, process loses frame
 - If actual rate too high, process gains frame



Working Sets and Page Fault Rates

31

- Direct relationship between working set of a process and its page fault rate
- Working set changes over time
- Peaks and valleys over time






Implementation examples of Paging & Segmentation in Intel Pentium

Example: The Intel Pentium

- Supports both segmentation and segmentation with paging
- CPU generates logical address
 - Given to segmentation unit
 - Which produces linear addresses
 - Linear address given to paging unit
 - Which generates physical address in main memory
 - Paging units form equivalent of MMU

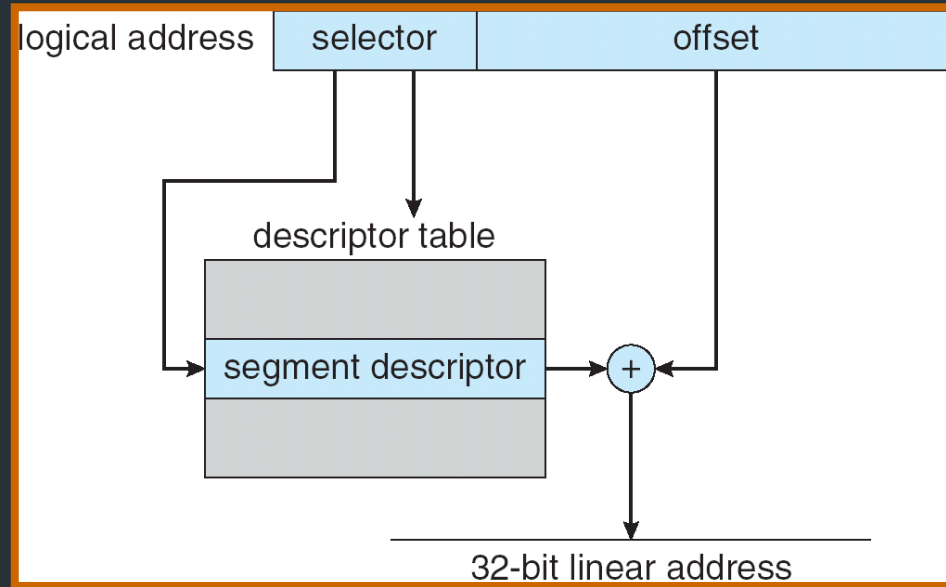


Logical to Physical Address Translation in Pentium

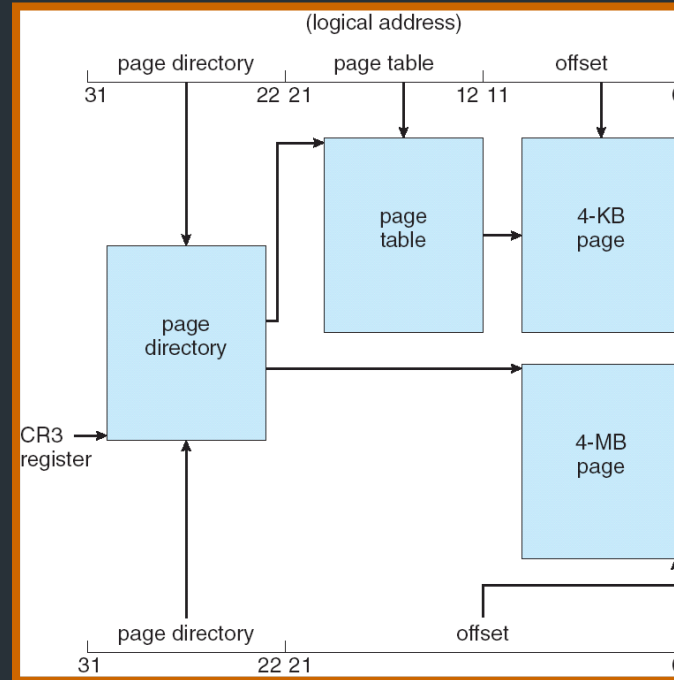


page number		page offset
p_1	p_2	d
10	10	12

Intel Pentium Segmentation



Pentium Paging Architecture





Any Questions ?

References

- <http://www.digilife.be/quickreferences/qrc/linux%20system%20call%20quick%20reference.pdf>
- <http://www.cs.nthu.edu.tw/~ychung/slides/CSC3150/Abraham-Silberschatz-Operating-System-Concepts---9th2012.12.pdf>
- "Applied Operating Systems Concepts", is based on the text "Operating System Concepts", 5/e (1998) by Abraham Silberschatz and Peter Baer Galvin.
- [Silberschatz, Gagne, Galvin: Operating System Concepts, 9th Edition](#)
- http://docs.cs.up.ac.za/programming/asm/derick_tut/syscalls.html

Thank You!