

File-System

File Concept

- A file is a collection of related information that is recorded on secondary storage
- Contiguous logical address space
- Types:
 - Data files
 - Program files
- File content is defined by its creator
- Many different types of information may be stored in a file - source or executable programs, numeric or text data, images, music, video, and so on.

File concept (contd...)

- A file has a certain defined structure, which depends on its type.
 - A **text file** is a sequence of characters organized into lines
 - A **source file** is a sequence of functions, each of which is further organized as declarations followed by executable statements
 - An **executable file** is a series of code sections that the loader can bring into memory and execute.

File Attributes

- **Name** – only information kept in human-readable form
- **Identifier** – unique tag (number) identifies file within file system
- **Type** – needed for systems that support different types
- **Location** – pointer to file location on device
- **Size** – current file size
- **Protection** – controls who can do reading, writing, executing
- **Time, date, and user identification** – data for protection, security, and usage monitoring

File Operations

- Creating a file
- Writing a file
- Reading a file
- Repositioning within a file - **seek**
- Deleting a file
- Truncating a file
- Renaming a file

Open Files

- The OS Keeps a table, called the **open-file table**, containing information about all open files.
- Several pieces of data are associated with an open files:
 - **File pointer**: pointer to last read/write location, per process that has the file open
 - **File-open count**: counter of number of times a file is open
 - **Disk location of the file**: cache of data access information
 - **Access rights**: per-process access mode information

Open File Locking

- Provided by some operating systems and file systems
 - Similar to reader-writer locks
 - **Shared lock** similar to reader lock – several processes can acquire concurrently
 - **Exclusive lock** similar to writer lock
- Operating systems may provide either **mandatory** or **advisory** file-locking mechanisms:
 - **Mandatory** – access is denied depending on locks held and requested
 - **Advisory** – processes can find status of locks and decide what to do

File Types – Name, Extension

| file type | usual extension | function |
|----------------|-----------------------------|--|
| executable | exe, com, bin or none | ready-to-run machine- language program |
| object | obj, o | compiled, machine language, not linked |
| source code | c, cc, java, pas, asm, a | source code in various languages |
| batch | bat, sh | commands to the command interpreter |
| text | txt, doc | textual data, documents |
| word processor | wp, tex, rtf, doc | various word-processor formats |
| library | lib, a, so, dll | libraries of routines for programmers |
| print or view | ps, pdf, jpg | ASCII or binary file in a format for printing or viewing |
| archive | arc, zip, tar | related files grouped into one file, sometimes com- pressed, for archiving or storage |
| multimedia | mpeg, mov, rm, mp3, avi | binary file containing audio or A/V information |

Access Methods

- The information in the file can be accessed in several ways
 - **Sequential Access** - Information in the file is processed in order, one record after the other
 - **Direct Access or Relative Access** - a file is made up of fixed-length **logical records** that allow programs to read and write records rapidly in no particular order
 - The direct-access method is based on a disk model of a file, since disks allow random access to any file block
 - **Indexed Access** – it involves in the construction of an index for the file.
 - The **index** contains pointers to the various blocks.
 - To find a record in the file, we first search the index and then use the pointer to access the file directly and to find the desired record.

Directories

- A directory contains information about a group of files
- Each entry in a directory contains the attributes of one file

| <i>File name</i> | <i>Type and size</i> | <i>Location info</i> | <i>Protection info</i> | <i>Open count</i> | <i>Lock</i> | <i>Flags</i> | <i>Misc info</i> |
|------------------|----------------------|----------------------|------------------------|-------------------|-------------|--------------|------------------|
| | | | | | | | |

- File name - Name of the file
- Type and size -The file's type and size
- Location info - Information about the file's location on a disk.
- Protection info - Information about which users are permitted to access this file and in what manner.

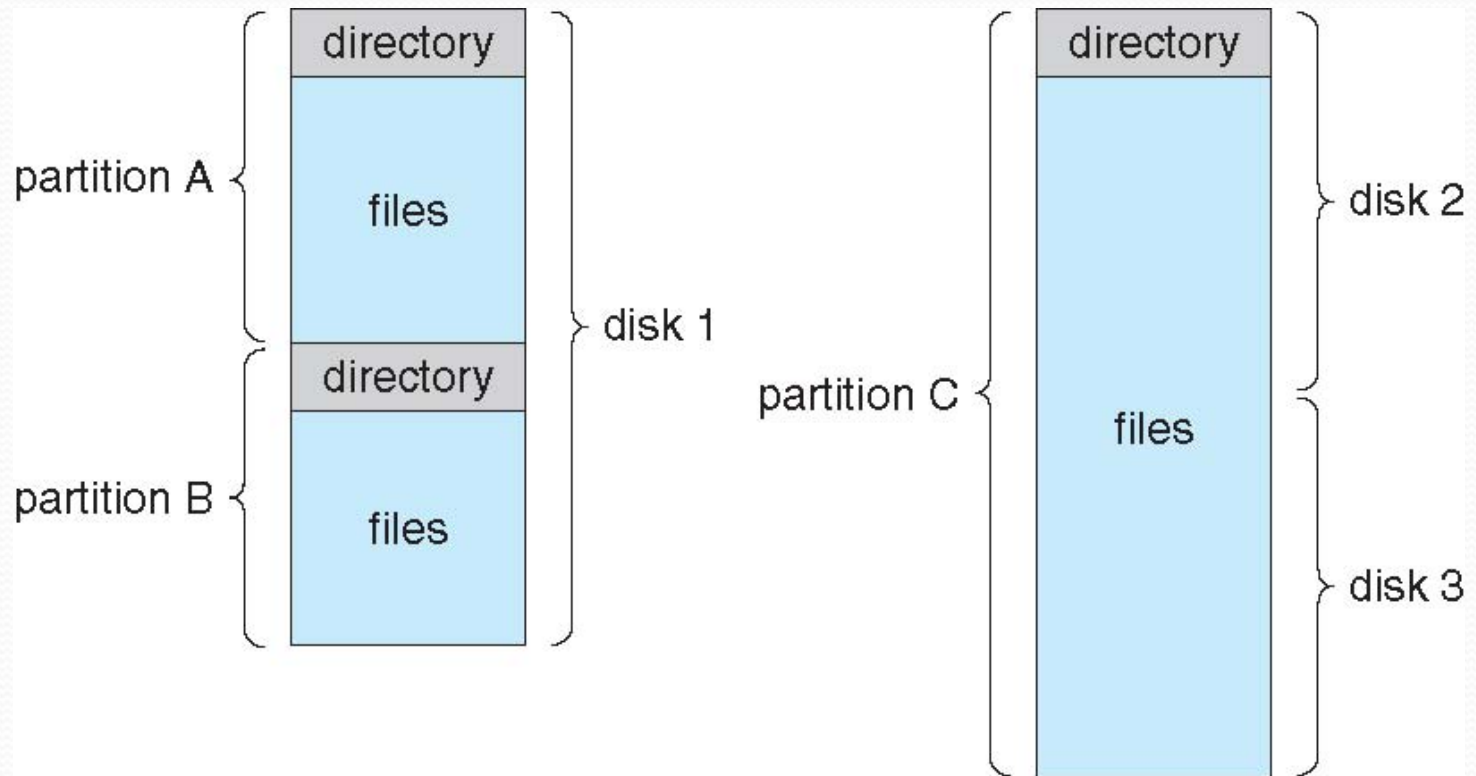
Directories (contd...)

- Open count - Number of processes currently accessing the file.
- Lock - Indicates whether a process is currently accessing the file in an exclusive manner.
- Flags - Information about the nature of the file—whether the file is a directory, a link, or a mounted file system.
- Misc info - Miscellaneous information like id of owner, date and time of creation, last use, and last modification.

Directory and Disk Structure

- Disk can be subdivided into **partitions**
- Disks or partitions can be **RAID** protected against failure
- Disk or partition can be used **raw** – without a file system, or **formatted** with a file system
- Partitions also known as minidisks, slices
- Entity containing file system known as a **volume**

A Typical File-system Organization



Directory

The directory can be viewed as a symbol table that translates file names into their directory entries

Operations performed on Directory

- Search for a file
- Create a file
- Delete a file
- List a directory
- Rename a file
- Traverse the file system


Directory Organization

The directory is organized logically to obtain

- Efficiency – locating a file quickly
- Naming – convenient to users
 - Two users can have same name for different files
 - The same file can have several different names
- Grouping – logical grouping of files by properties, (e.g., all Java programs, all games, ...)

Single-Level Directory

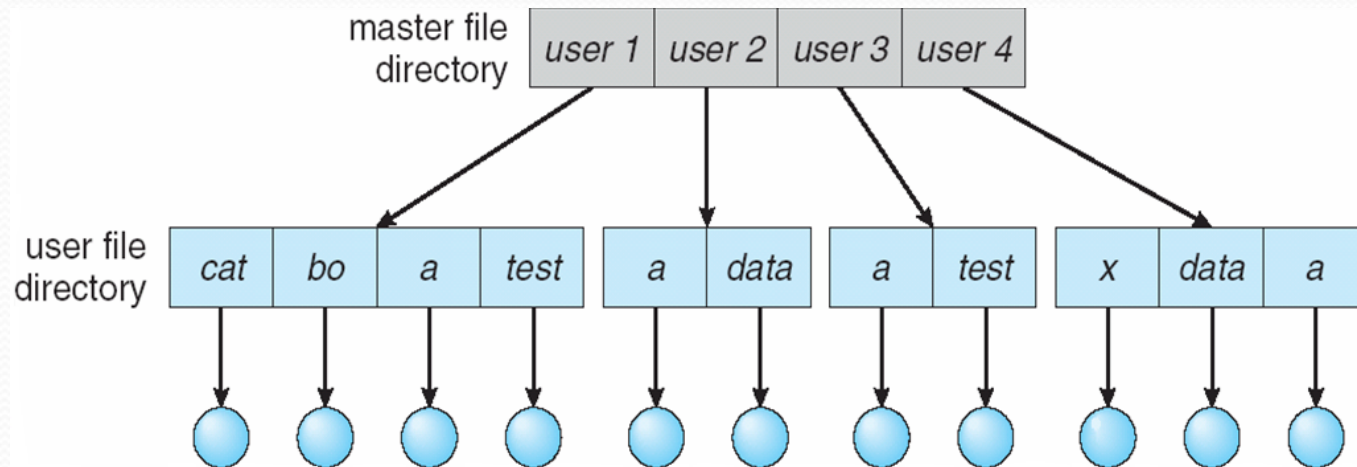
- A single directory for all users



- Naming problem
- Grouping problem

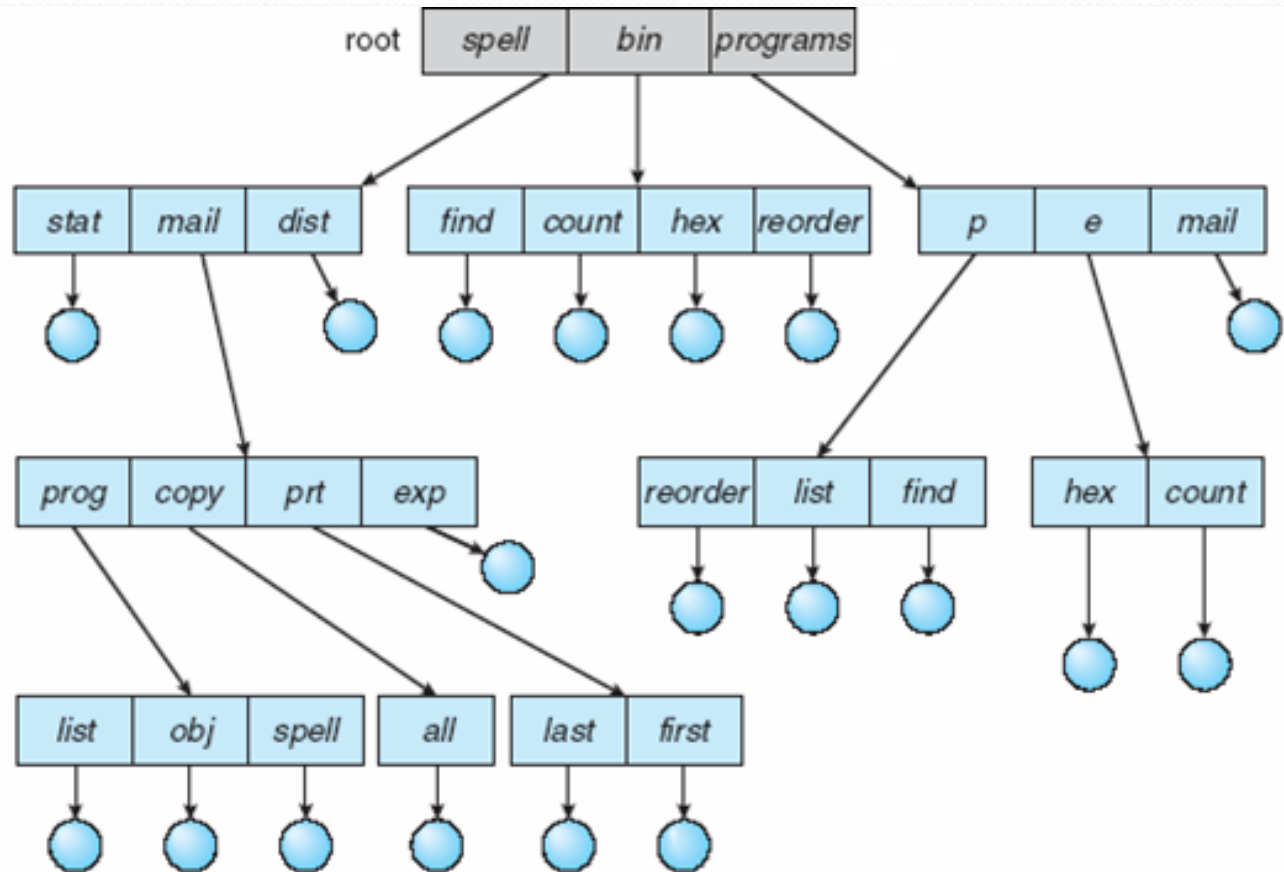
Two-Level Directory

- Separate directory for each user



- Path name
- Can have the same file name for different user
- Efficient searching
- No grouping capability

Tree-Structured Directories



Tree-Structured Directories (Cont.)

Advantage:

- Efficient searching
- Grouping Capability

Disadvantage :

- Sharing of file – not possible

Tree-Structured Directories (Cont)

Path names can be of two types:

- An **absolute path name** begins at the root and follows a path down to the specified file, giving the directory names on the path.
- A **relative path name** defines a path from the current directory.

An interesting policy decision in a tree-structured directory concerns how to handle the deletion of a directory.

If a directory is empty, its entry in the directory that contains it can simply be deleted.

Acyclic-Graph Directories

- Shared files and subdirectories can be implemented in several ways.
- A common way is to create a new directory entry called a link.
- A **link** is effectively a pointer to another file or subdirectory
- An acyclic-graph directory structure is more flexible than a simple tree structure, but it is also more complex.
- A file may now have multiple absolute path names.
- Another problem involves deletion

File System Mounting

- Each file system is constituted on a *logical disk*, i.e., on a partition of a disk.
- Files contained in a file system can be accessed only when the file system is *mounted*
- The mount operation is what “connects” the file system to the system’s directory structure
- An unmount operation disconnects a file system.
- The operating system is given the name of the device and the **mount point**—the location within the file structure where the file system is to be attached
- Typically, a mount point is an empty directory
- A unmounted file system is mounted at a **mount point**

Mounting of a File System

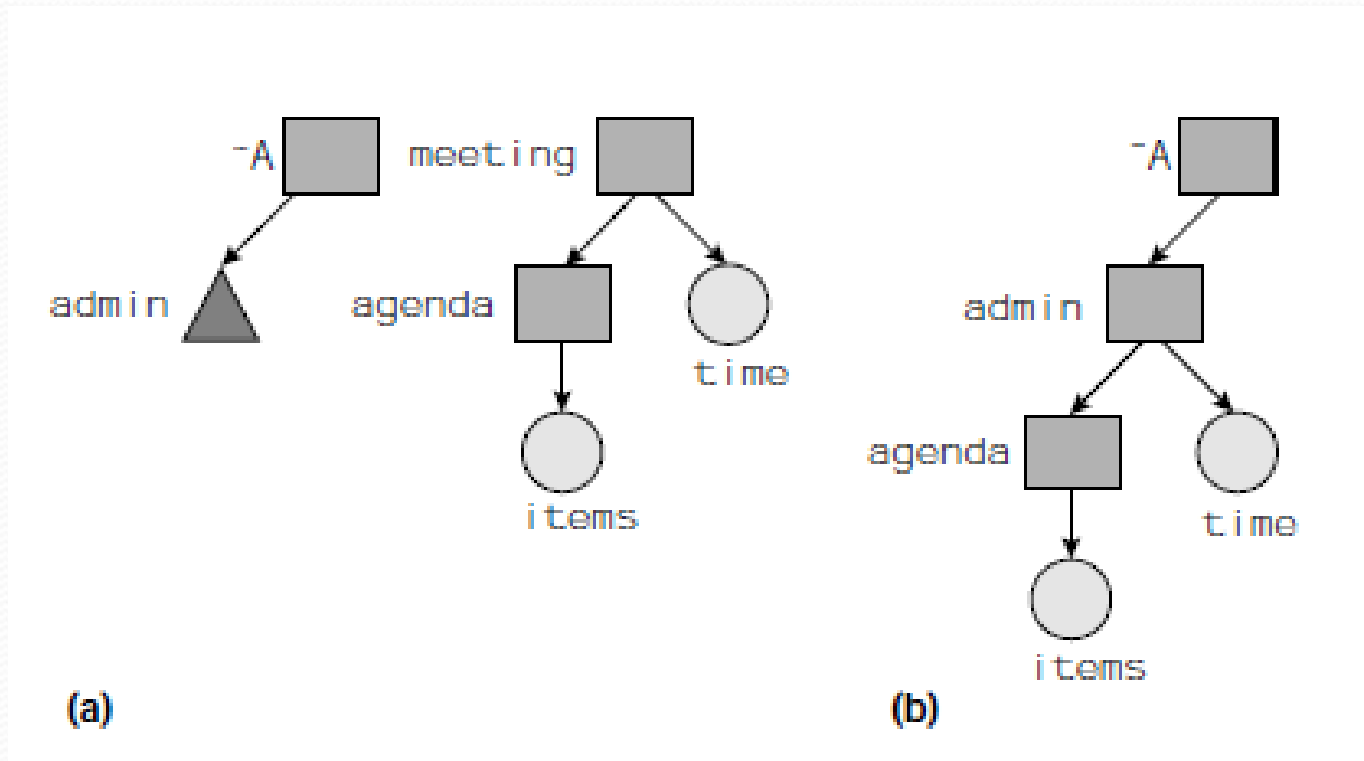


Fig (a) ~A/admin is a mount point in a directory structure, and meeting is the root directory of another file system.

Fig (b) shows the effect of the command mount (meeting, ~A/admin).

Directory Implementation

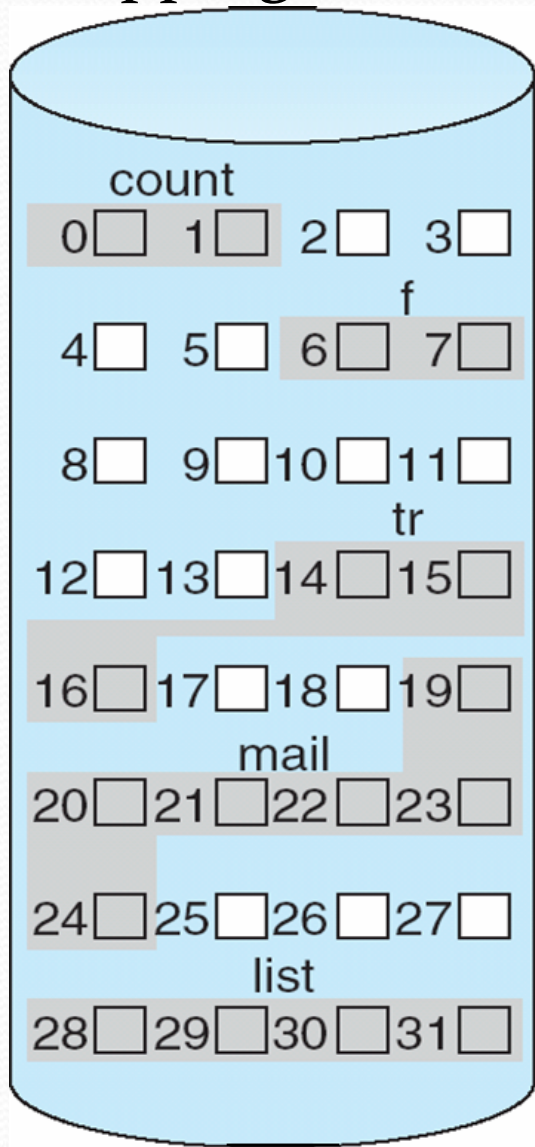
- **Linear list** of file names with pointer to the data blocks
 - Simple to program
 - Time-consuming to execute
 - Linear search time
 - Could keep ordered alphabetically via linked list or use B+ tree
- **Hash Table** – linear list with hash data structure
 - Decreases directory search time
 - **Collisions** – situations where two file names hash to the same location
 - Only good if entries are fixed size, or use chained-overflow method

Allocation Methods - Contiguous

- An allocation method refers to how disk blocks are allocated for files:
- **Contiguous allocation** – each file occupies set of contiguous blocks
 - Best performance in most cases
 - Simple – only starting location (block #) and length (number of blocks) are required
 - Problems include finding space for file, knowing file size, external fragmentation, need for **compaction off-line (downtime)** or **on-line**

Contiguous Allocation

- Mapping from logical to physical




directory

| file | start | length |
|-------|-------|--------|
| count | 0 | 2 |
| tr | 14 | 3 |
| mail | 19 | 6 |
| list | 28 | 4 |
| f | 6 | 2 |

Allocation Methods - Linked

- **Linked allocation** – each file a linked list of blocks
 - File ends at nil pointer
 - No external fragmentation
 - Each block contains pointer to next block
 - No compaction, external fragmentation
 - Free space management system called when new block needed
 - Improve efficiency by clustering blocks into groups but increases internal fragmentation
 - Locating a block can take many I/Os and disk seeks

Linked Allocation



Allocation Methods – Linked (Cont.)

- FAT (File Allocation Table) variation
 - Beginning of volume has table, indexed by block number
 - Much like a linked list, but faster on disk and cacheable
 - New block allocation simple


File-Allocation Table

directory entry

| | | |
|------|-----|-----|
| test | ... | 217 |
|------|-----|-----|

name

start block



no. of disk blocks


-1

FAT

Allocation Methods - Indexed

- Indexed allocation

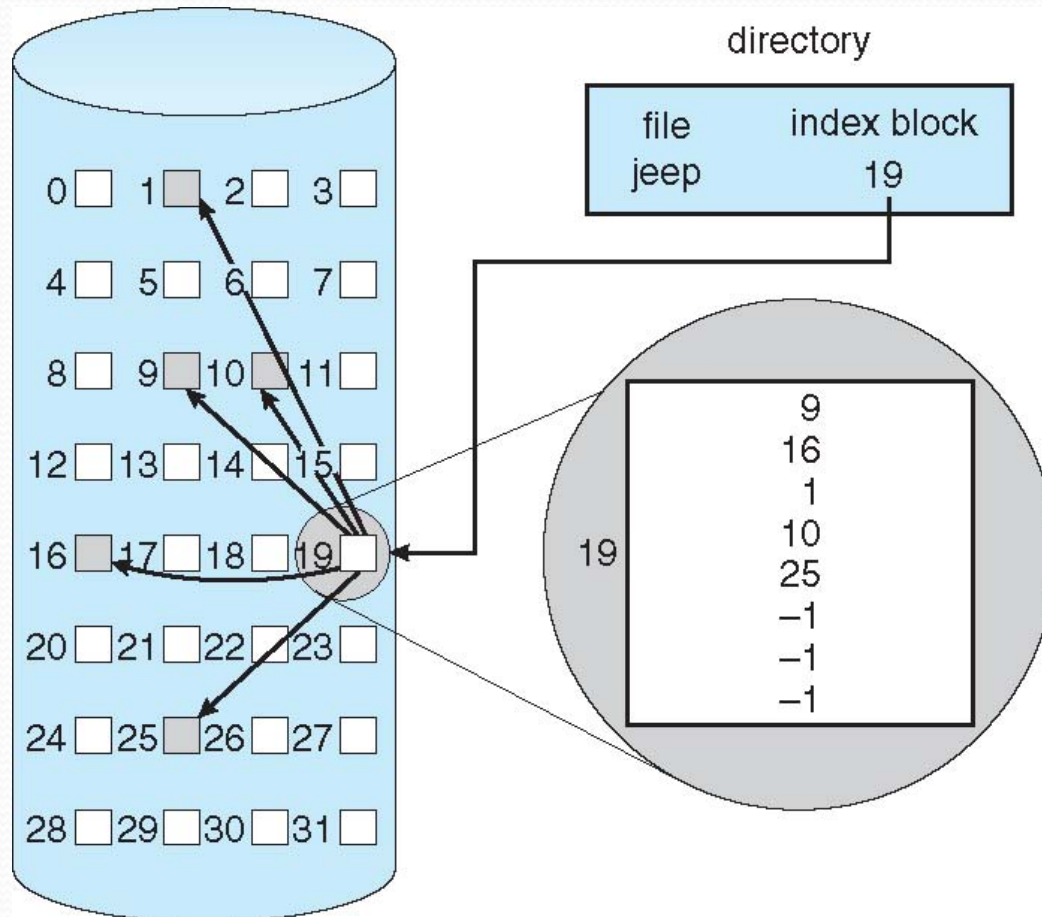
- Each file has its own **index block**(s) of pointers to its data blocks



- Logical view

index table

Example of Indexed Allocation



Indexed Allocation (Cont.)

- Need index table
- Random access
- Dynamic access without external fragmentation, but have overhead of index block
- We need a block for index table



Distributed File System

Terms in DFS Context

- **Service** - software entity running on one or more machines and providing a particular type of function to clients
- **Server** is the service software running on a single machine
- **Client** is a process that can invoke a service using a set of operations that form its **client interface**
- File system provides file services to clients
- A client interface for a file service is formed by a set of primitive file operations, such as create a file, delete a file, read from a file, and write to a file.

DFS

- A DFS is a file system whose clients, servers, and storage devices are dispersed among the machines of a distributed system
- Instead of a single centralized data repository, the system frequently has multiple and independent storage devices
- The distinctive features of a DFS are the multiplicity and autonomy of clients and servers in the system
- The most important performance measure of a DFS is the amount of time needed to satisfy service requests and its transparency

Naming and Transparency

- **Naming** is a mapping between logical and physical objects
- Usually, a user refers to a file by a textual name
- It is mapped to a lower-level numerical identifier that in turn is mapped to disk blocks
- This multilevel mapping provides users with an abstraction of a file that hides the details of how and where on the disk the file is stored.
- In a transparent DFS, a new dimension is added to the abstraction: that of hiding where in the network the file is located
- In this abstraction, both the existence of multiple copies and their locations are hidden

Naming Structures

We need to differentiate two related notions regarding name mappings in a DFS:

- 1. Location transparency.** The name of a file does not reveal any hint of the file's physical storage location
- 2. Location independence.** The name of a file does not need to be changed when the file's physical storage location changes

Naming Schemes

There are three main approaches to naming schemes in a DFS.

1. In the simplest approach, a file is identified by some combination of its host name and local name, which guarantees a unique system-wide name

- The Internet URL system also uses this approach. This naming scheme is neither location transparent nor location independent

Naming Schemes (contd...)

2. . The second approach was popularized by Sun's network file system

- NFS provides a means to attach remote directories to local directories, thus giving the appearance of a coherent directory tree
- **Automount** feature allowed mounts to be done on demand based on a table of mount points and file-structure names

Naming Schemes (contd...)

3. We can achieve total integration of the component file systems by using the third approach.
 - Here, a single global name structure spans all the files in the system
 - In practice, it is difficult to attain

Remote File Access

- Requests for accesses to files are delivered to the server, the server machine performs the accesses, and their results are forwarded back to the user.
- One of the most common ways of Implementing remote service is the RPC paradigm
- Caching –is used to ensure reasonable performance of a remote-service mechanism
- The main goal of caching is to reduce both network traffic and disk I/O

Basic Caching Scheme

- The idea is to retain recently accessed disk blocks in the cache, so that repeated accesses to the same information can be handled locally, without additional network traffic
- Files are still identified with one master copy residing at the server machine, but copies (or parts) of the file are scattered in different caches
- When a cached copy is modified, the changes need to be reflected on the master copy to preserve the relevant consistency semantics

Cache Location

- Where should the cached data be stored—on disk or in main memory?
- Disk caches have one clear advantage over main-memory caches: they are reliable
- Modifications to cached data are lost in a crash if the cache is kept in volatile memory
- Moreover, if the cached data are kept on disk, they are still there during recovery, and there is no need to fetch them again

Cache Location

Main-memory caches have several advantages of their own, however:

- Main-memory caches permit workstations to be diskless
- Data can be accessed more quickly from a cache in main memory than from one on a disk
- The server caches will be in main memory regardless of where user caches are located

Cache-Update Policy

- The policy used to write modified data blocks back to the server's master copy has a critical effect on the system's performance and reliability.
1. **write-through policy** - The simplest policy is to write data through to disk as soon as they are placed in any cache.
- Advantage - reliability: little information is lost when a client system crashes.
 - However, this policy requires each write access to wait until the information is sent to the server, so it causes poor write performance.

Cache-Update Policy

2. **Delayed-write policy**, also known as **write-back caching**, where we delay updates to the master copy.

- Modifications are written to the cache and then are written through to the server at a later time.
- Advantages:
 - Because writes are made to the cache, write accesses complete much more quickly.
 - Only the last update needs to be written
- Disadvantages :
 - reliability problems, since unwritten data are lost whenever a user machine crashes.

Cache-Update Policy

3. write-on-close policy - write data back to the server when the file is closed.

- For files that are open for long periods and are modified frequently, however, the performance advantages of this policy over delayed write with more frequent flushing are apparent.

Consistency

- If the client machine determines that its cached data are out of date, it must cache an up-to-date copy of the data before allowing further accesses.
- There are two approaches to verifying the validity of cached data:
 1. **Client-initiated approach.** The client initiates a validity check, in which it contacts the server and checks whether the local data are consistent with the master copy.
- The frequency of the validity checking is the crux of this approach

Consistency

2. Server-initiated approach. The server records, for each client, the files (or parts of files) that it caches.

- When the server detects a potential inconsistency, it must react.



Any Queries?