




**VELLORE INSTITUTE OF TECHNOLOGY**  
**WINTER SEMESTER 2021**

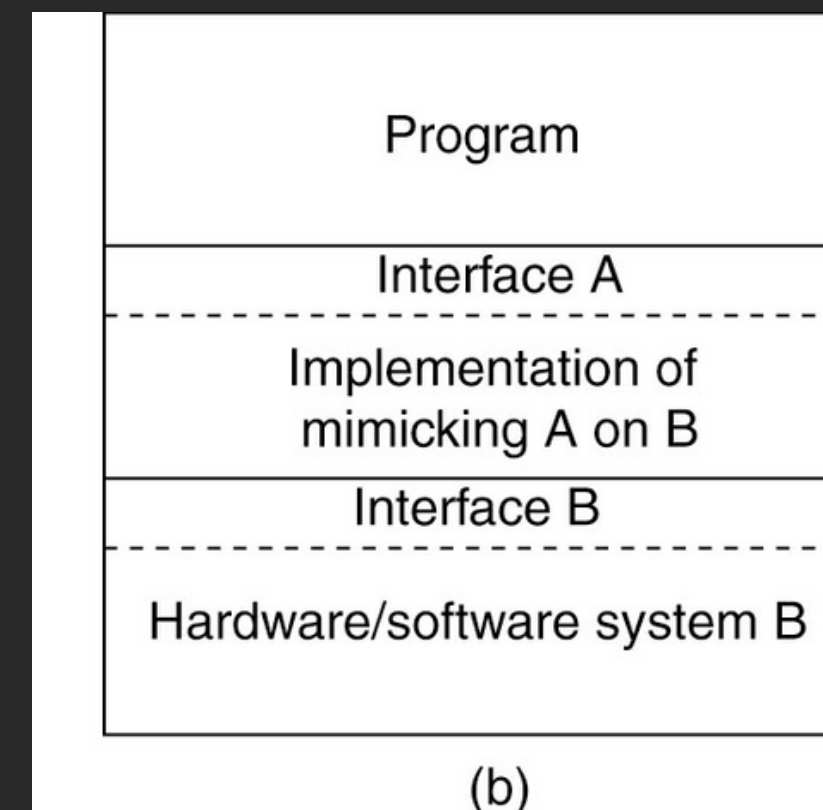
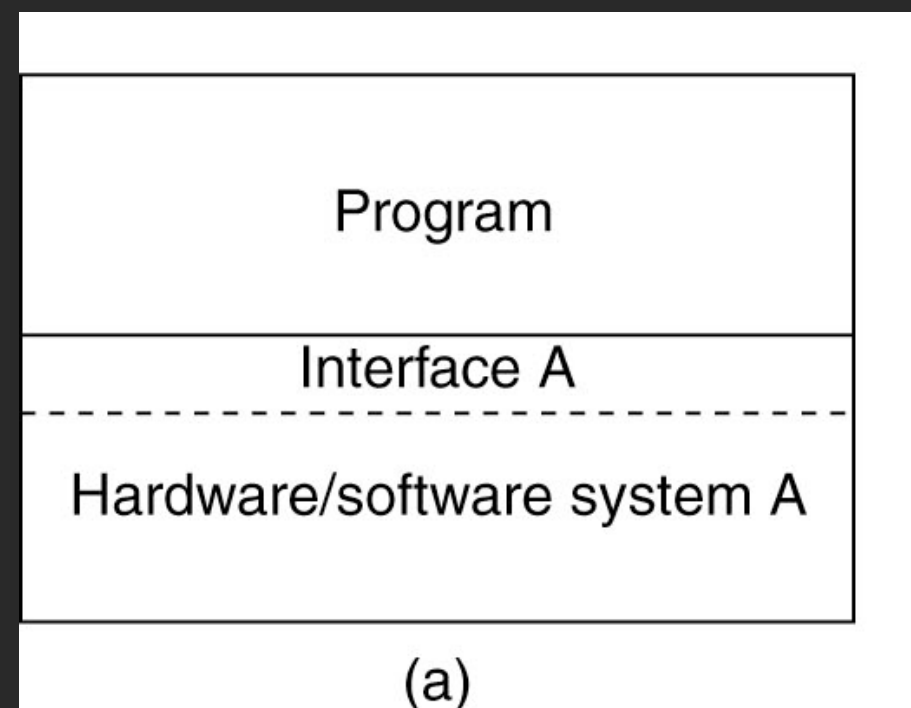
# **CSE2005 OPERATING SYSTEMS**

# **VIRTUALIZATION**



# VIRTUALIZATION

- **Virtualization: extend or replace an existing interface to mimic the behavior of another system.**
  - Introduced in 1970s: run legacy software on newer mainframe hardware
- **Handle platform diversity by running apps in VMs**
  - Portability and flexibility





# VIRTUALIZATION MACHINES

- **Virtualization technology enables a single PC or server to simultaneously run multiple operating systems or multiple sessions of a single OS**
- **A machine with virtualization software can host numerous applications, including those that run on different operating systems, on a single platform**
- **The host operating system can support a number of virtual machines, each of which has the characteristics of a particular OS**
- **The solution that enables virtualization is a virtual machine monitor (VMM), or hypervisor**



# APPROACHES TO VIRTUALIZATION

- ▶▶ it is configured with some number of processors, some amount of RAM, storage resources, and connectivity through the network ports
- ▶▶ once the VM is created it can be powered on like a physical server, loaded with an operating system and software solutions, and utilized in the manner of a physical server
- ▶▶ unlike a physical server, this virtual server only sees the resources it has been configured with, not all of the resources of the physical host itself

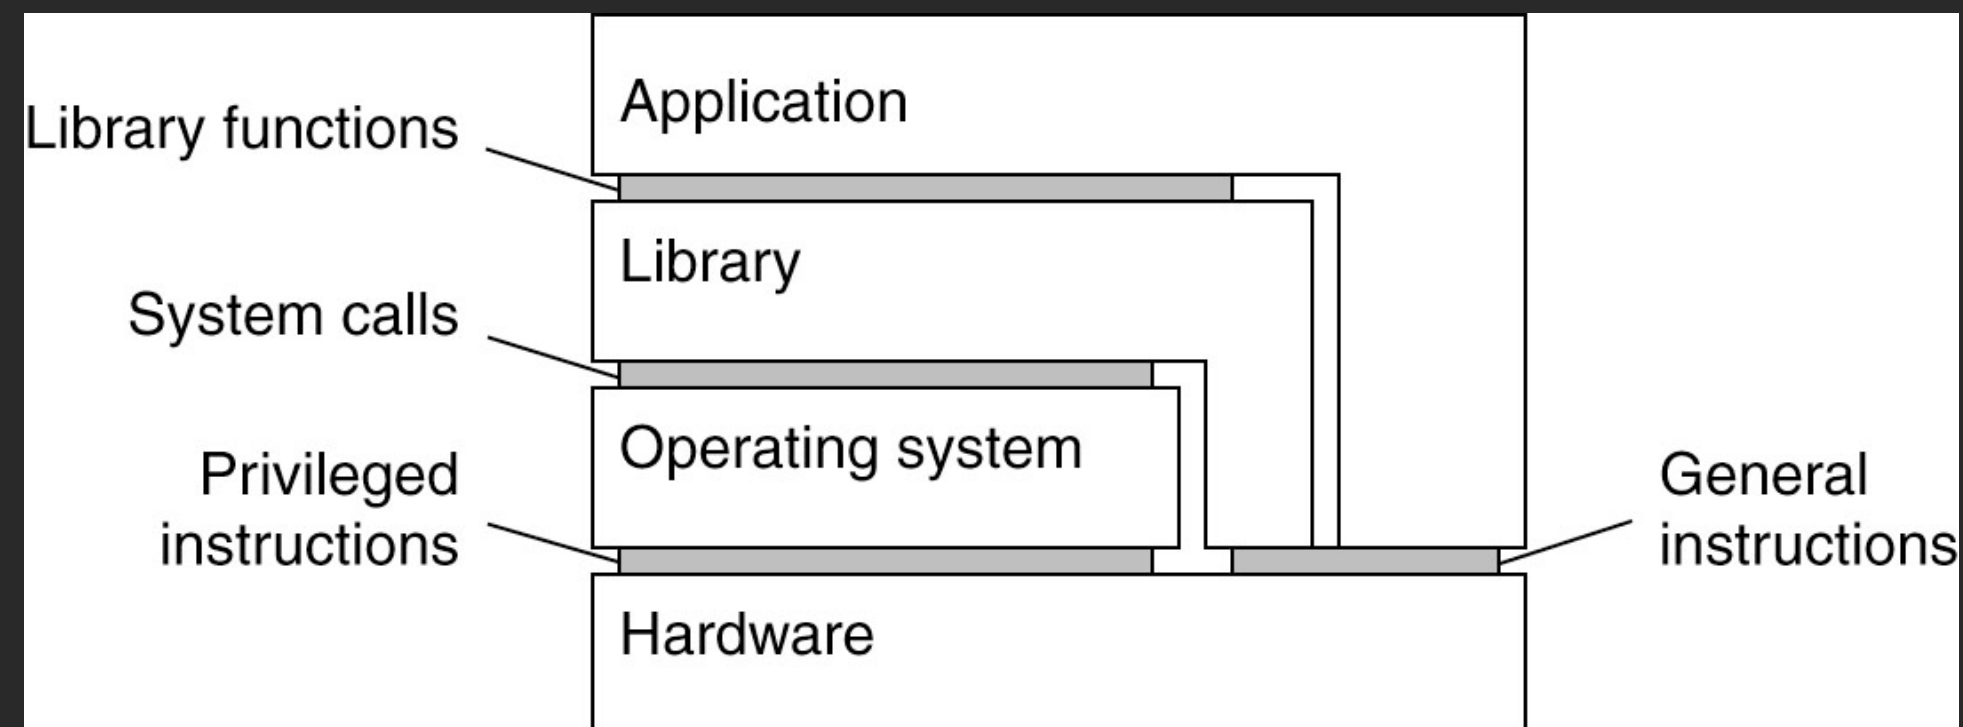
# **VIRTUAL MACHINES ARE MADE UP OF FILES:**

- ▶▶ **configuration file describes the attributes of the virtual machine**
- ▶▶ **it contains the server definition, how many virtual processors (vCPUs) are allocated to this virtual machine, how much RAM is allocated, which I/O devices the VM has access to, how many network interface cards (NICs) are in the virtual server, and more**
- ▶▶ **it also describes the storage that the VM can access**
- ▶▶ **when a virtual machine is powered on, or instantiated, additional files are created for logging, for memory paging, and other functions**
- ▶▶ **since VMs are already files, copying them produces not only a backup of the data but also a copy of the entire server, including the operating system, applications, and the hardware configuration itself**



# TYPES OF INTERFACES

- Different types of interfaces
  - Assembly instructions
  - System calls
  - APIs
- Depending on what is replaced /mimiced, we obtain different forms of virtualization





# TYPES OF VIRTUALIZATION

- **Emulation**
  - VM emulates/simulates complete hardware
  - Unmodified guest OS for a different PC can be run
- **Bochs, VirtualPC for Mac, QEMU**
- **Full/native Virtualization**
  - VM simulates “enough” hardware to allow an unmodified guest OS to be run in isolation
- **Same hardware CPU**
  - IBM VM family, VMWare Workstation, Parallels, VirtualBox





# TYPES OF VIRTUALIZATION

- **Para-virtualization**
  - VM does not simulate hardware
  - Use special API that a modified guest OS must use
  - Hypercalls trapped by the Hypervisor and serviced
  - Xen, VMWare ESX Server
- **OS-level virtualization**
  - OS allows multiple secure virtual servers to be run
  - Guest OS is the same as the host OS, but appears isolated
- **apps see an isolated OS**
  - Solaris Containers, BSD Jails, Linux Vserver, Linux containers, Docker
- **Application level virtualization**
  - Application is gives its own copy of components that are not shared
- **(E.g., own registry files, global objects) - VE prevents conflicts**
  - JVM, Rosetta on Mac (also emulation), WINE







# TYPES OF HYPERVISORS

- Hypervisor/VMM: virtualization layer
  - resource management, isolation, scheduling, ...
- Type 1: hypervisor runs on “bare metal”
- Type 2: hypervisor runs on a host OS
  - Guest OS runs inside hypervisor
- Both VM types act like real hardware





# HOW VIRTUALIZATION WORKS?

- CPU supports kernel and user mode (ring0, ring3)
  - Set of instructions that can only be executed in kernel mode
- I/O, change MMU settings etc -- sensitive instructions
  - Privileged instructions: cause a trap when executed in user mode
- Result: type 1 virtualization feasible if sensitive instruction subset of privileged instructions
- Intel 386: ignores sensitive instructions in user mode
  - Can not support type 1 virtualization
- Recent Intel/AMD CPUs have hardware support
  - Intel VT, AMD SVM
- Create containers where a VM and guest can run
- Hypervisor uses hardware bitmap to specify which inst should trap
- Sensitive inst in guest traps to hypervisor



# TYPE I HYPERVISOR

- **Unmodified OS is running in user mode (or ring 1)**
  - **But it thinks it is running in kernel mode (virtual kernel mode)**
  - **privileged instructions trap; sensitive inst-> use VT to trap**
  - **Hypervisor is the “real kernel”**
- **Upon trap, executes privileged operations**
- **Or emulates what the hardware would do**



# TYPE 2 HYPERVISOR

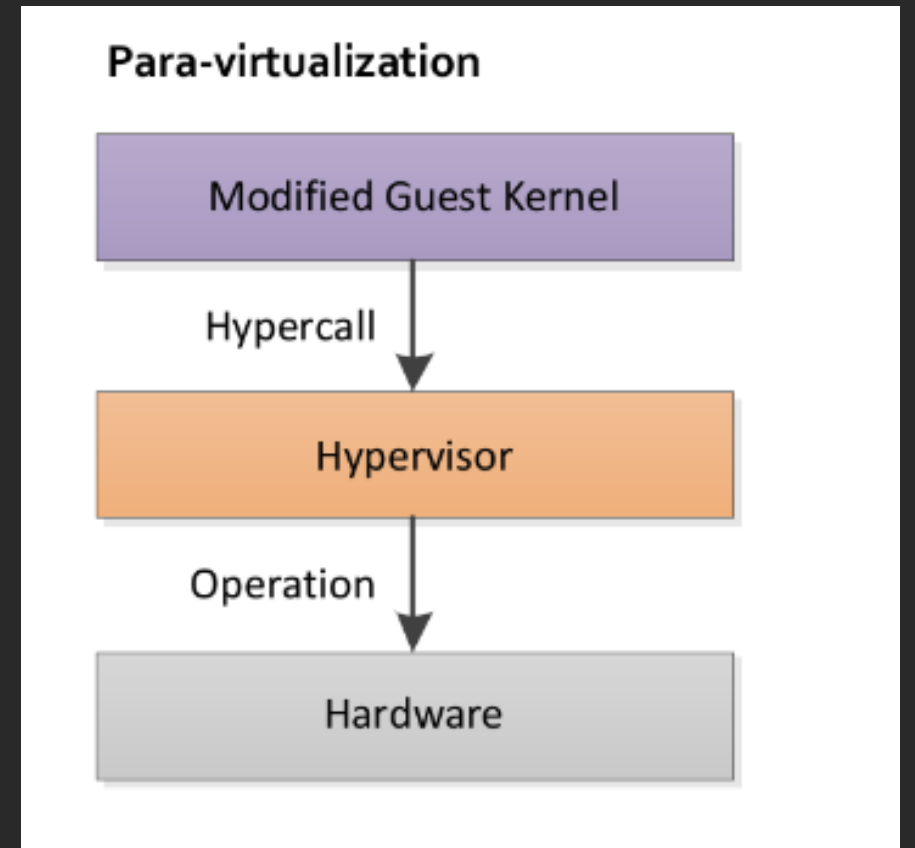
- **VMWare example**
  - Upon loading program: scans code for basic blocks
  - If sensitive instructions, replace by Vmware procedure
- **Binary translation**
  - Cache modified basic block in VMWare cache
- **Execute; load next basic block etc.**
- **Type 2 hypervisors work without VT support**
  - Sensitive instructions replaced by procedures that emulate them.





# PARAVIRTUALIZATION

- Both type 1 and 2 hypervisors work on unmodified OS
- Paravirtualization: modify OS kernel to replace all sensitive instructions with hypercalls
  - OS behaves like a user program making system calls
  - Hypervisor executes the privileged operation invoked by hypercall.





# MEMORY VIRTUALIZATION

- OS manages page tables
  - Create new pagetable is sensitive -> traps to hypervisor
- hypervisor manages multiple OS
  - Need a second shadow page table
  - OS: VM virtual pages to VM's physical pages
  - Hypervisor maps to actual page in shadow page table
  - Two level mapping
  - Need to catch changes to page table (not privileged)
- Change PT to read-only - page fault
- Paravirtualized - use hypercalls to inform





# I/O VIRTUALIZATION

- Each guest OS thinks it “owns” the disk
- Hypervisor creates “virtual disks”
  - Large empty files on the physical disk that appear as “disks” to the guest OS
- Hypervisor converts block # to file offset for I/O
  - DMA need physical addresses
- Hypervisor needs to translate
- NIC Virtualization





# VIRTUAL APPLIANCES & MULTI-CORE

- **Virtual appliance: pre-configured VM with OS/ apps pre-installed**
  - Just download and run (no need to install/configure)
  - Software distribution using appliances
- **Multi-core CPUs**
  - Run multiple VMs on multi-core systems
  - Each VM assigned one or more vCPU
  - Mapping from vCPUs to physical CPUs
- **Today: Virtual appliances have evolved into docker containers**







# USE OF VIRTUALIZATION TODAY

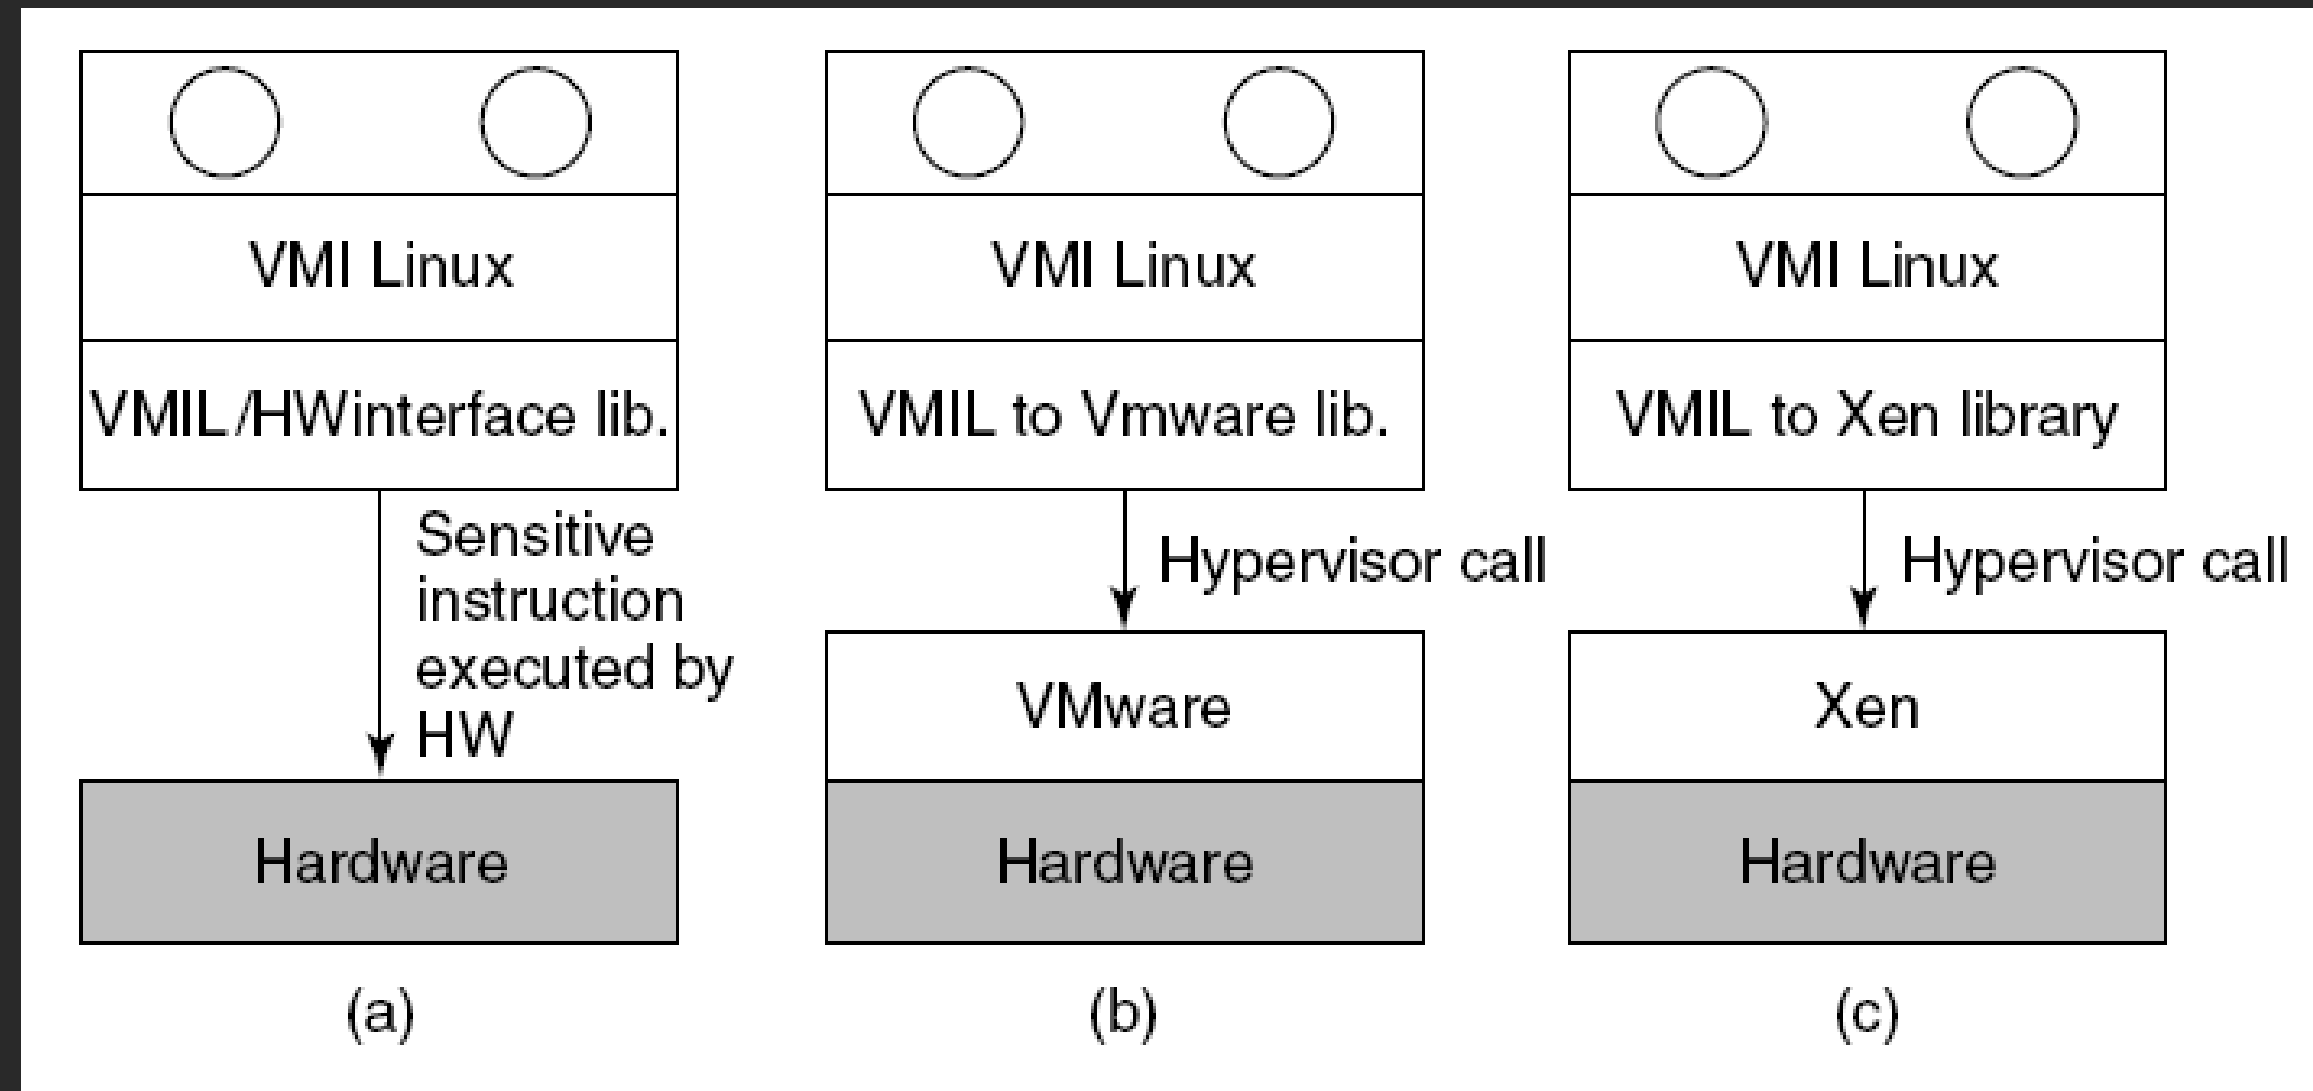
- **Data centers:**
  - **server consolidation: pack multiple virtual servers onto a smaller number of physical server**
- **saves hardware costs, power and cooling costs**
- **Cloud computing: rent virtual servers**
  - **cloud provider controls physical machines and mapping of virtual servers to physical hosts**
  - **User gets root access on virtual server**
- **Desktop computing:**
  - **Multi-platform software development**
  - **Testing machines**
  - **Run apps from another platform**





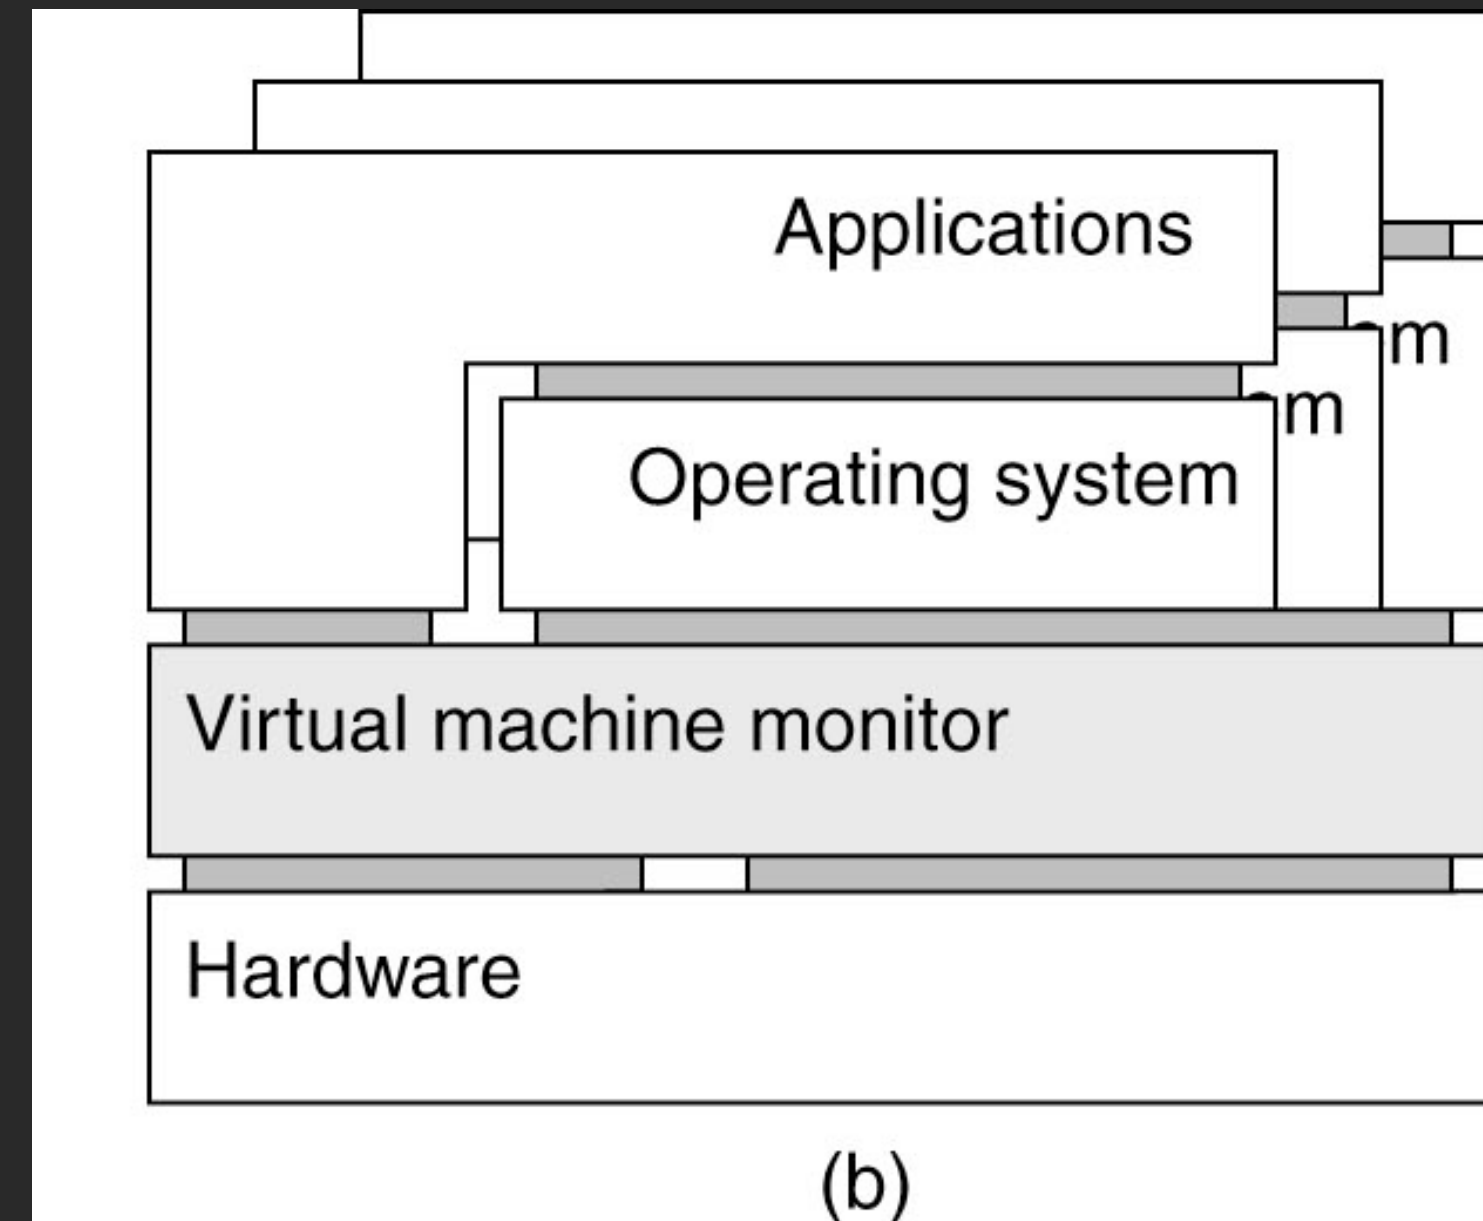
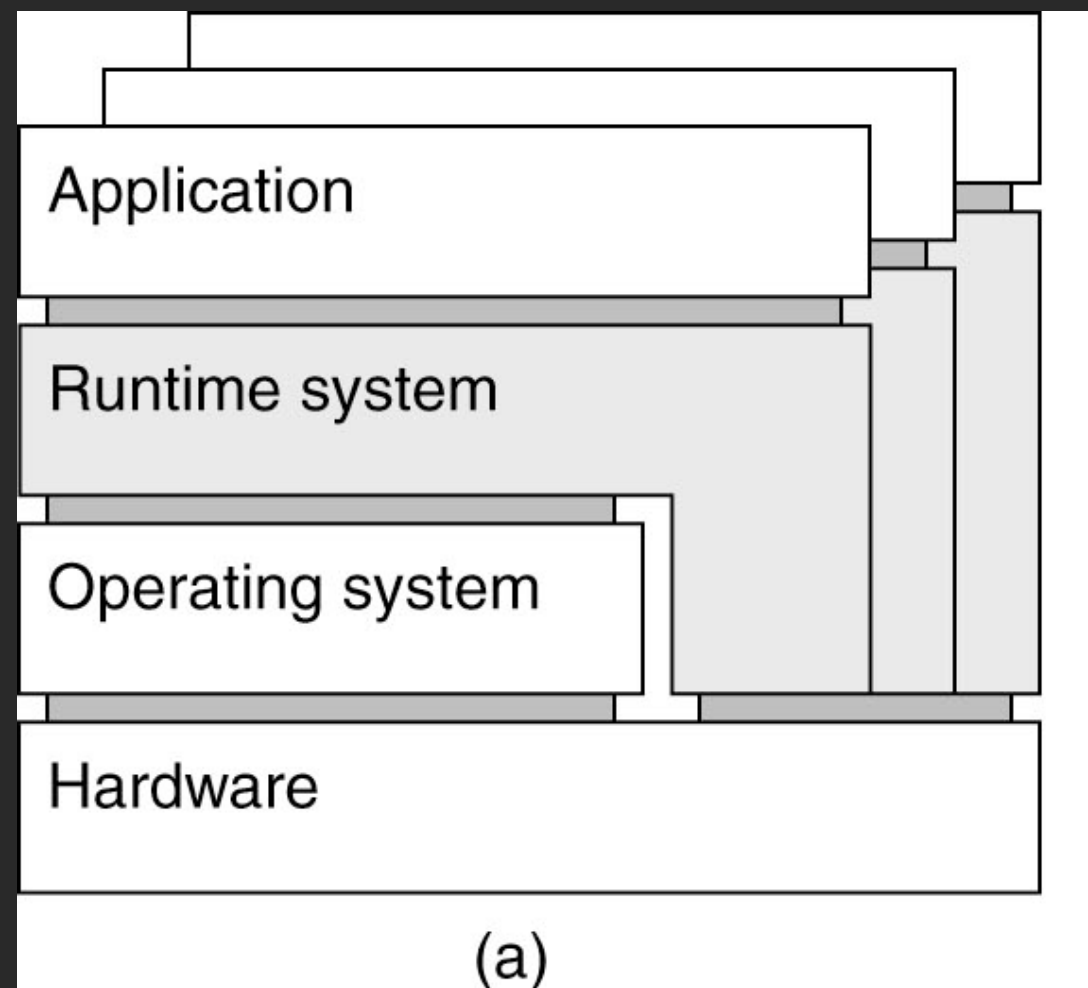
# VIRTUAL MACHINE INTERFACE

- Standardize the VM interface so kernel can run on bare hardware or any hypervisor



# EXAMPLES

- Application-level virtualization: “process virtual machine”
- VMM /hypervisor



# JAVA VM

- The goal of a Java Virtual Machine (JVM) is to provide a runtime space for a set of Java code to run on any operating system staged on any hardware platform without needing to make code changes to accommodate the different operating systems or hardware
- The JVM can support multiple threads
- Promises “Write Once, Run Anywhere”
- The JVM is described as being an abstract computing machine consisting of:
  - an instruction set
  - a program counter register
  - a stack to hold variables and results
  - a heap for runtime data and garbage collection
  - a method area for code and constants



# LINUX VSERVER

- **Linux VServer is an open-source, fast, lightweight approach to implementing virtual machines on a Linux server**
- **Only a single copy of the Linux kernel is involved**
- **VServer consists of a relatively modest modification to the kernel plus a small set of OS userland tools**
- **The VServer Linux kernel supports a number of separate virtual servers**
- **The kernel manages all system resources and tasks, including process scheduling, memory, disk space, and processor time**



thank  
you